

# **Scaling resources for sustainable provision of URLLC services**

**Static and dynamic optimization**

**Pablo Serrano Yáñez-Mingot**

**Brescia, October 10, 2025**

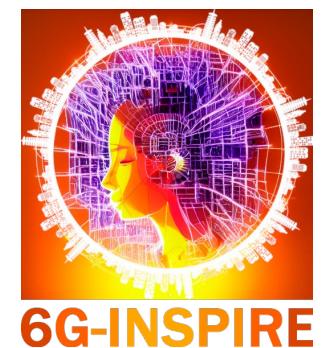
# Acknowledgements

This presentation is partly supported by

The Project TUCAN6-CM (TEC-2024/COM-460), funded by CM (ORDEN 5696/2024)



The Project 6GINSPIRE PID2022-137329OB-C42, funded by MCIN/AEI/10.13039/501100011033/



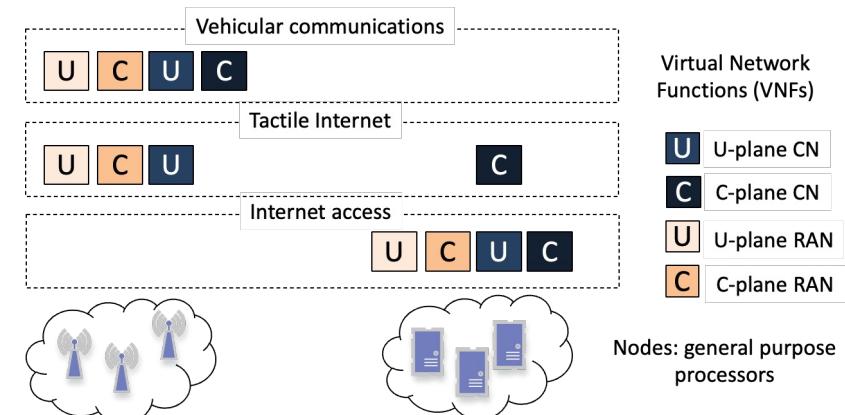
The SNS JU EU's HE research and innovation programme under Grant Agreement No. 101192035 (AMAZING-6G)



# Context

## *Network Softwarization*

- Network slicing: efficient provision of multiple heterogeneous services on the same infrastructure [1]
- Virtualization (softwarization) of network functions
  - Dynamically allocate and share resources (e.g., Nuberu [2])
  - Migrate between servers without disruption (e.g., ACHO [3])

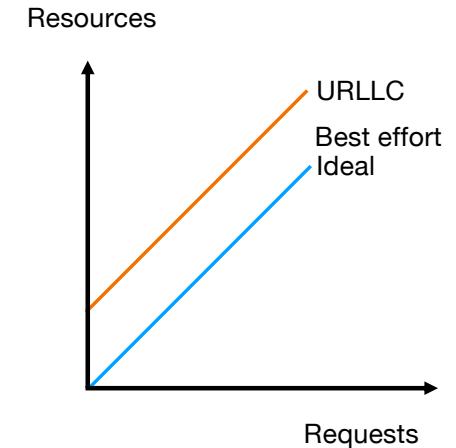


1. P. Serrano, «Tutorial: A Primer on 5G Network Slicing: Concepts, Algorithms, and Practice», IEEE CAMAD 2018, Barcelona,
2. G. Garcia-Aviles et al. «Nuberu: Reliable RAN Virtualization in Shared Platforms», ACM Mobicom '21
3. G. Garcia-Aviles et al. «ACHO: A Framework for Flexible Re-Orchestration of Virtual Network Functions», Computer Networks, 2020

# Efficiency with a URLLC service

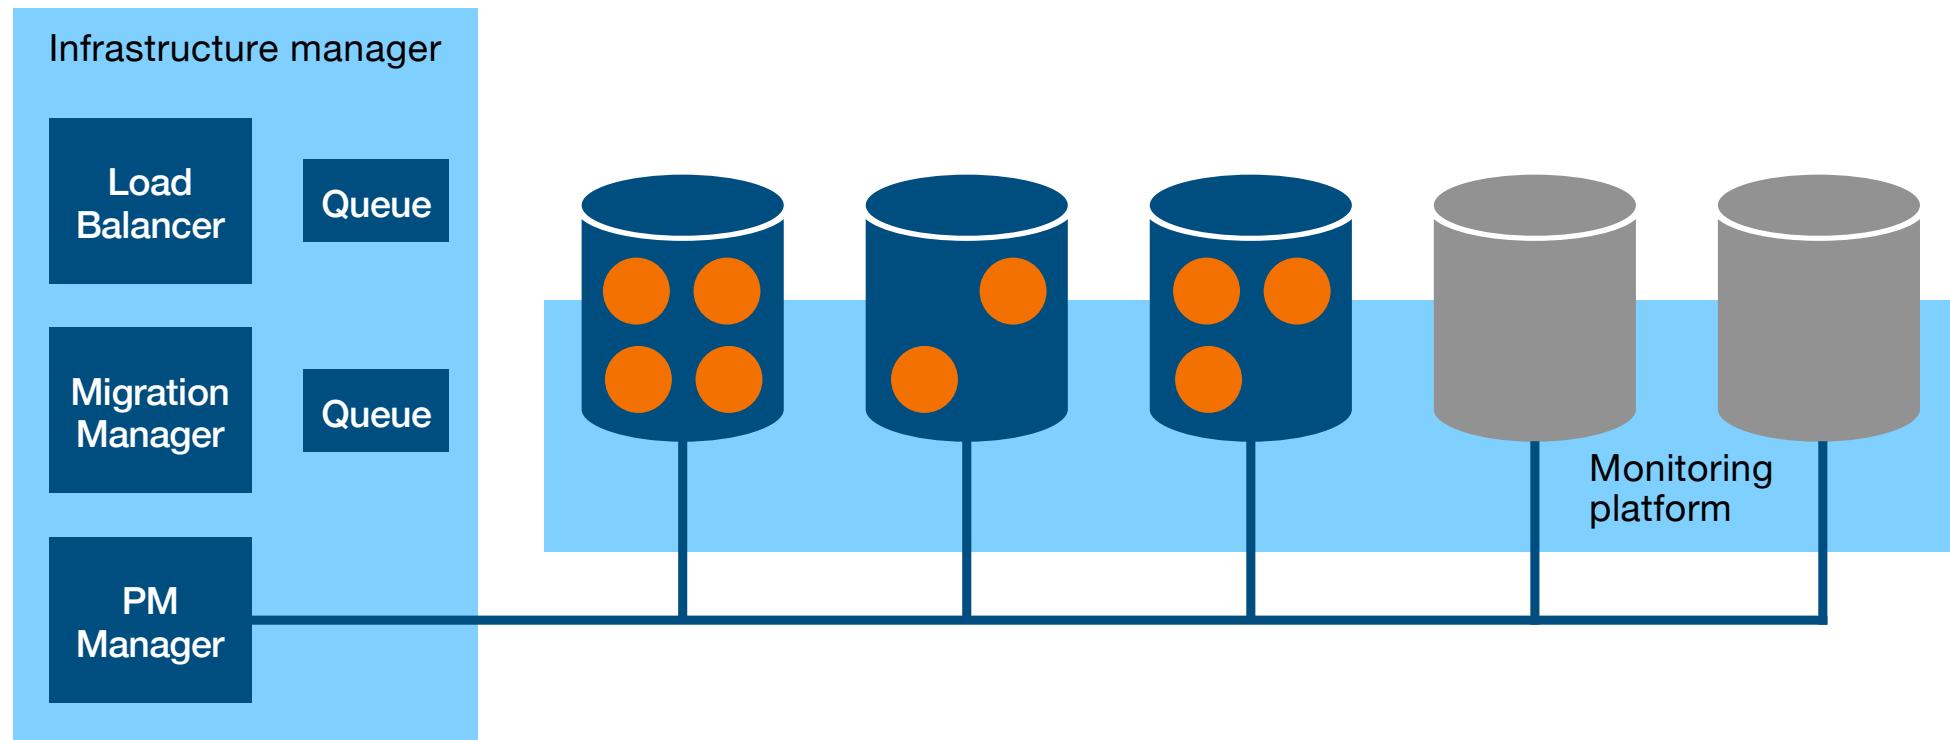
## Very high reliability

- Efficiency: allocate strictly necessary resources
- Obstacle 1: Resource bootstrapping is not instantaneous
- Bursts of requests: they may have to wait
- Obstacle 2: Life spans are finite [1]
  - Relocate tasks from servers that are about to fail
- Best effort vs. URLLC [2] (five 9s)

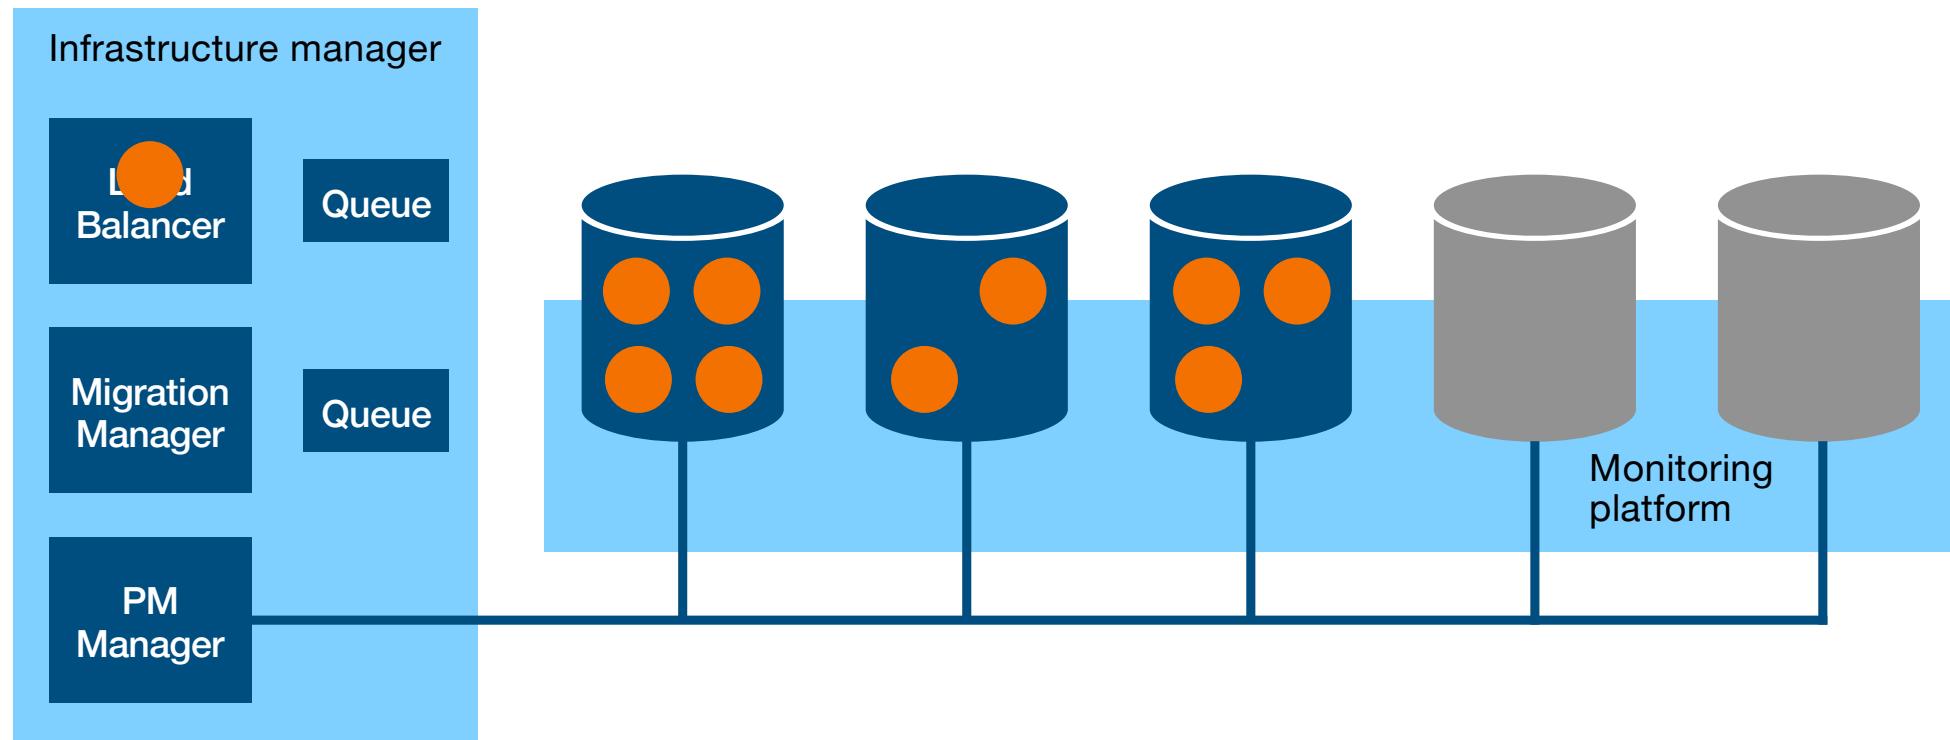


1. Fung Po Tso et al. «The Glasgow Raspberry Pi cloud: A scale model for cloud computing infrastructures». IEEE 33rd International Conference on Distributed Computing Systems Workshops, 2013.
2. W. Nakimuli et al. «Deployment and Evaluation of an Industry 4.0 Use Case over 5G», IEEE Communications Magazine, July 2021

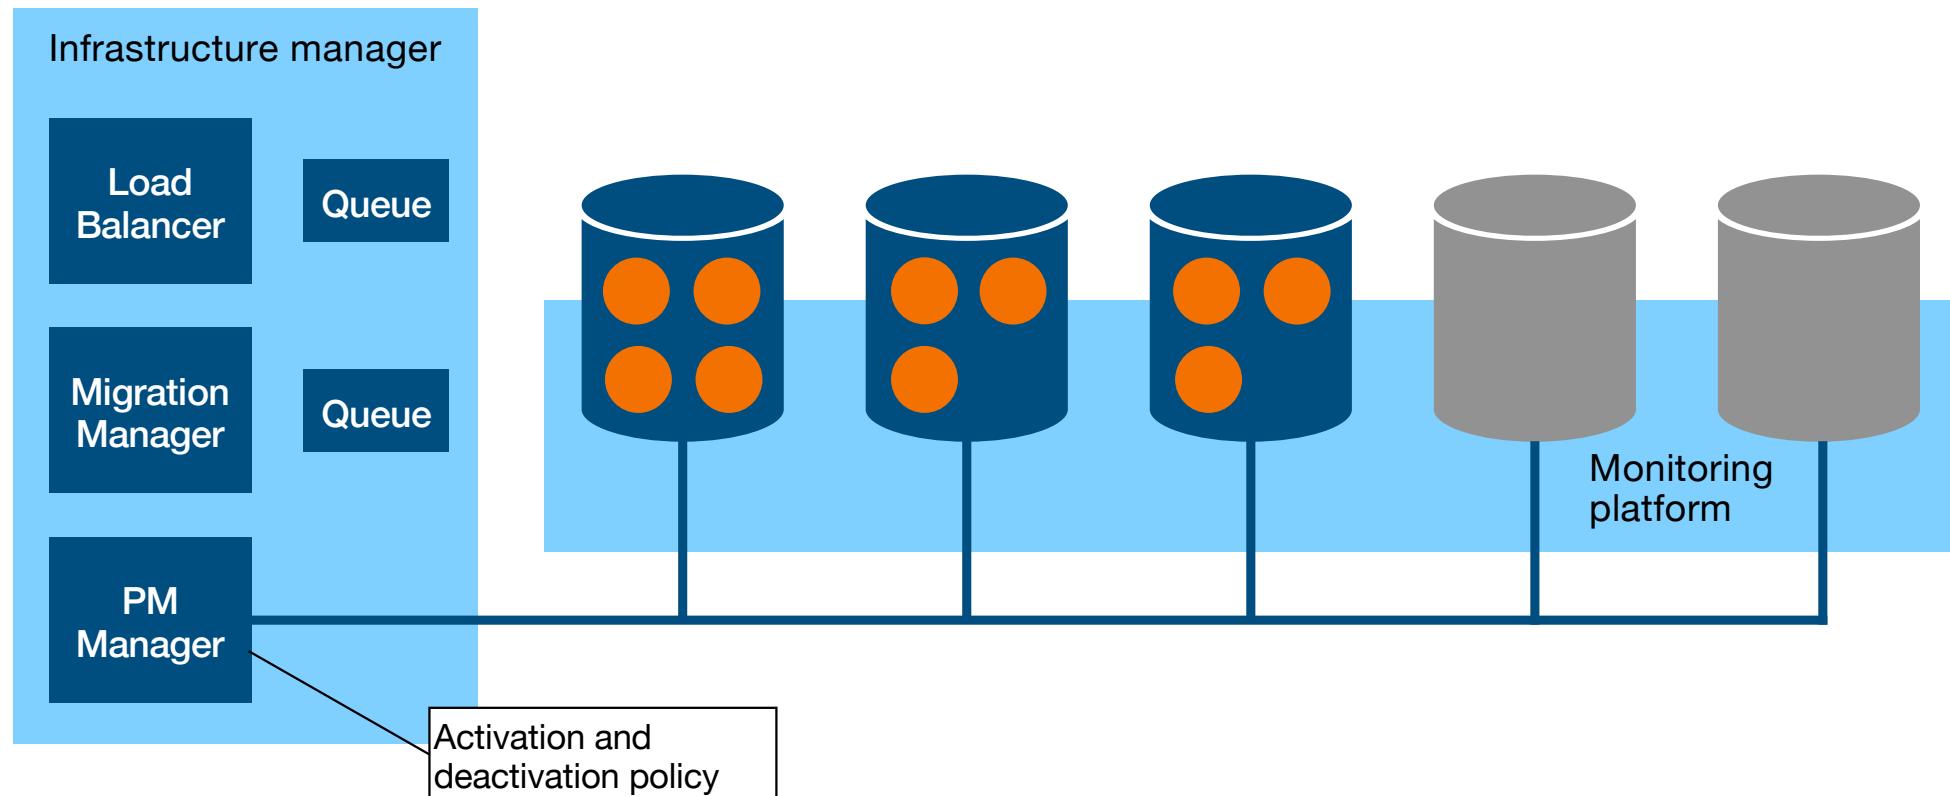
# Example



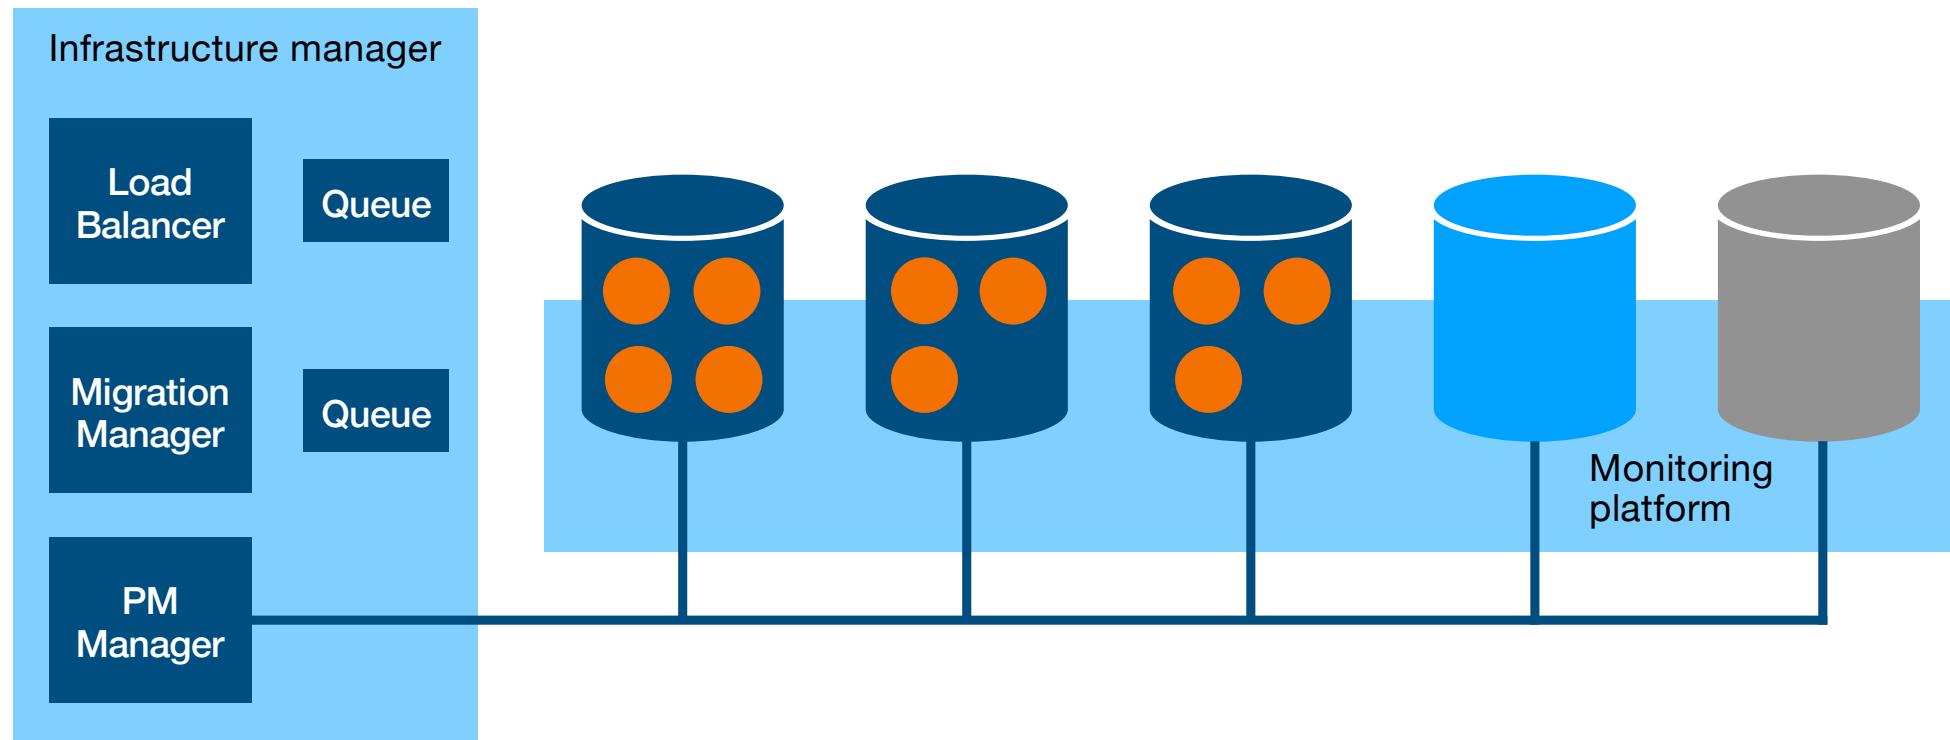
# Example



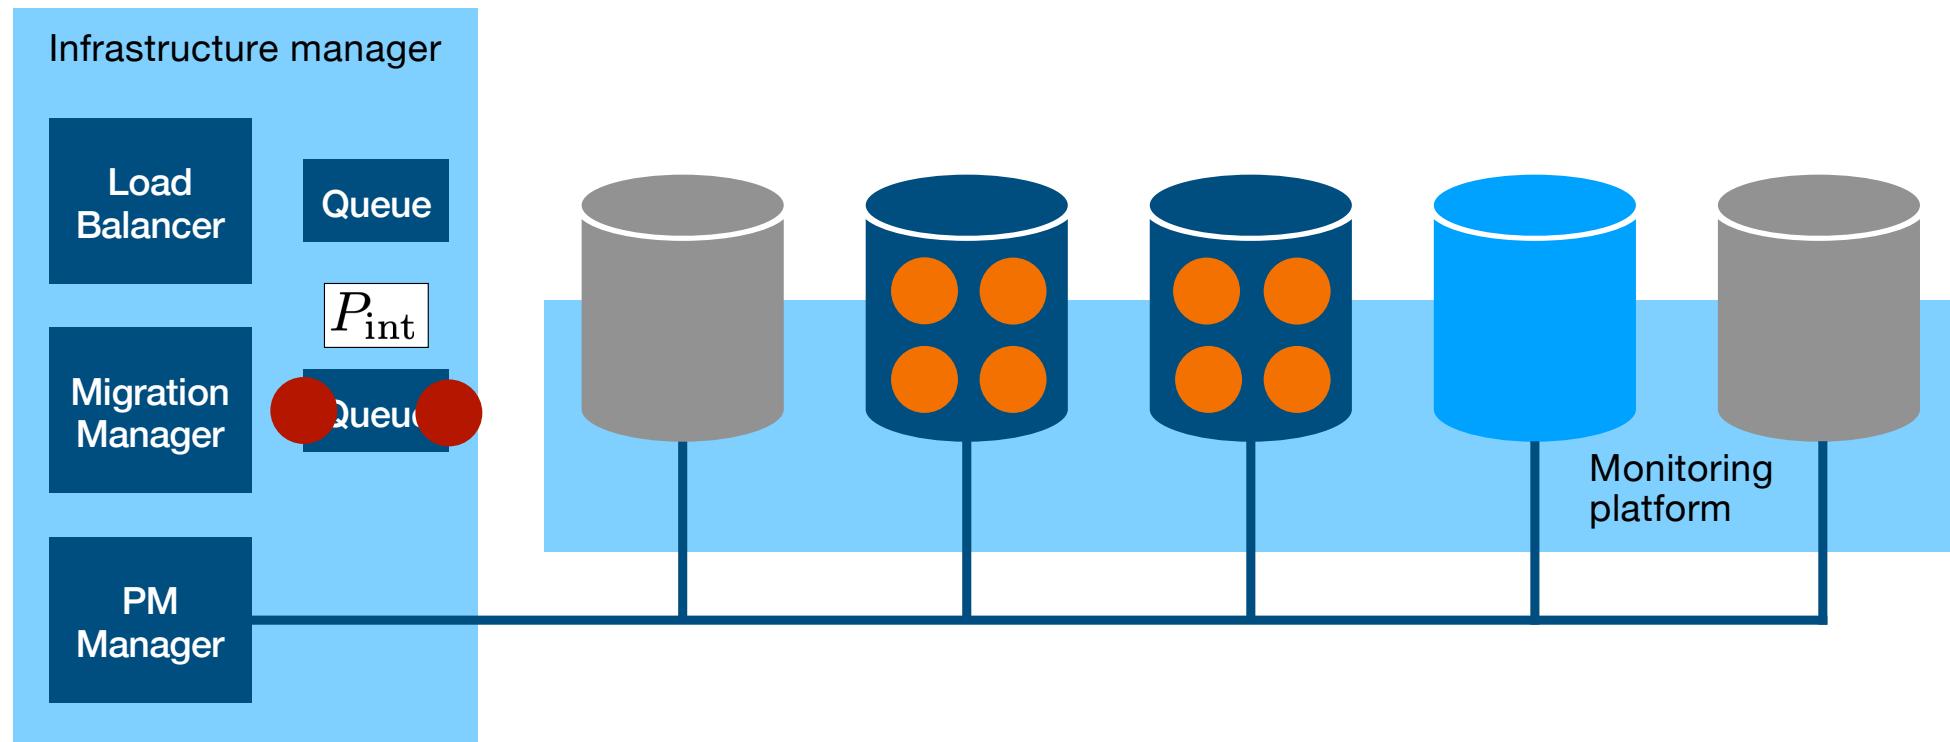
# Example



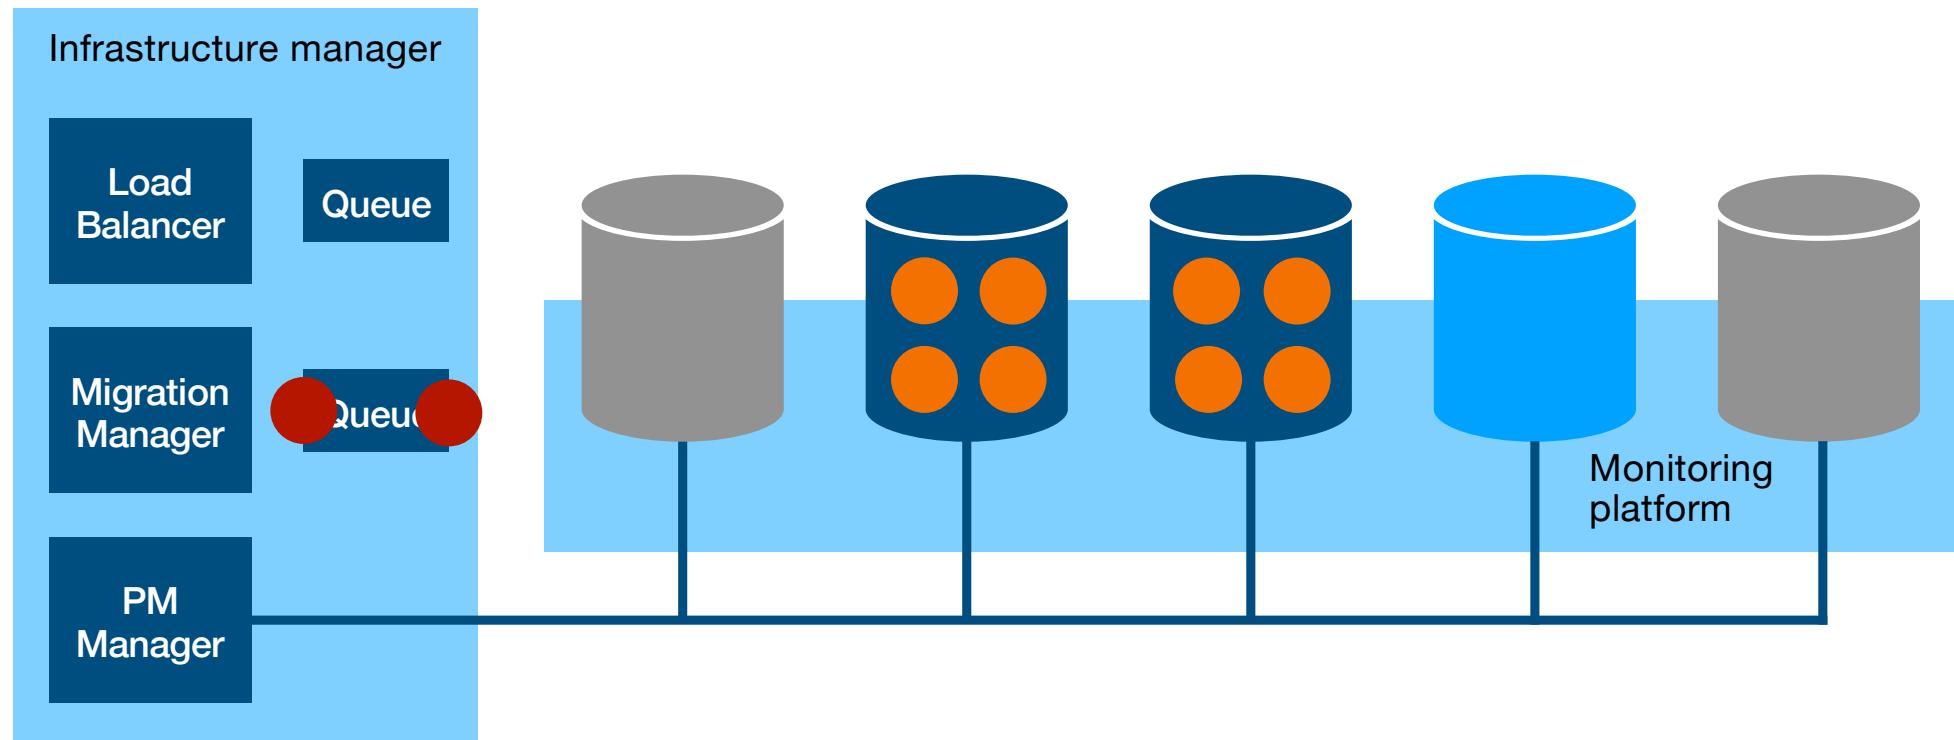
# Example



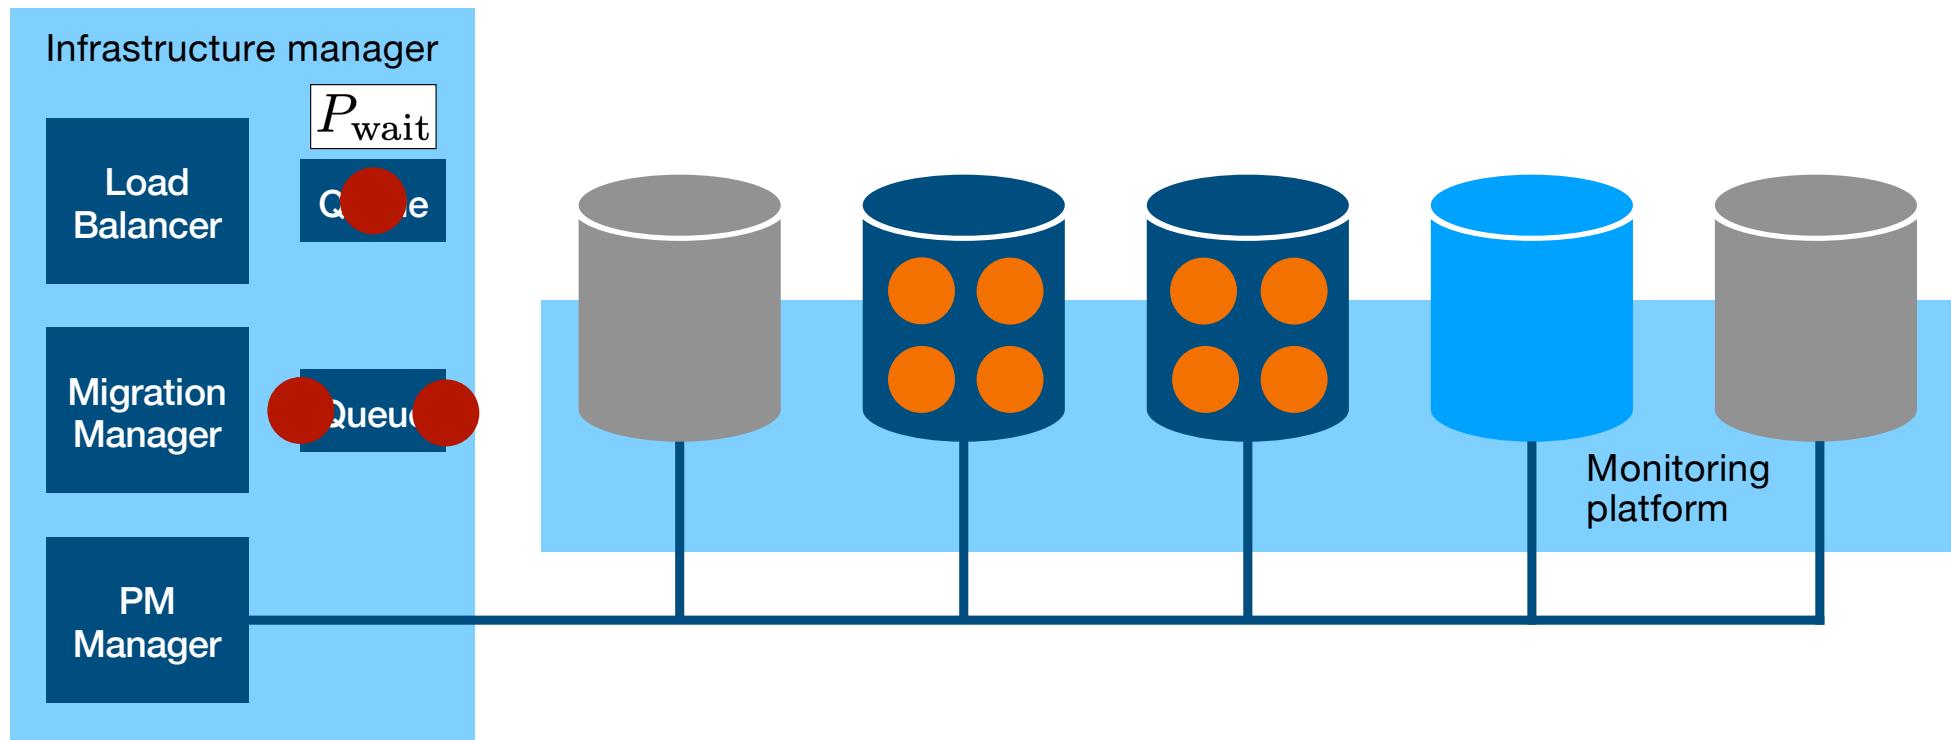
# Example



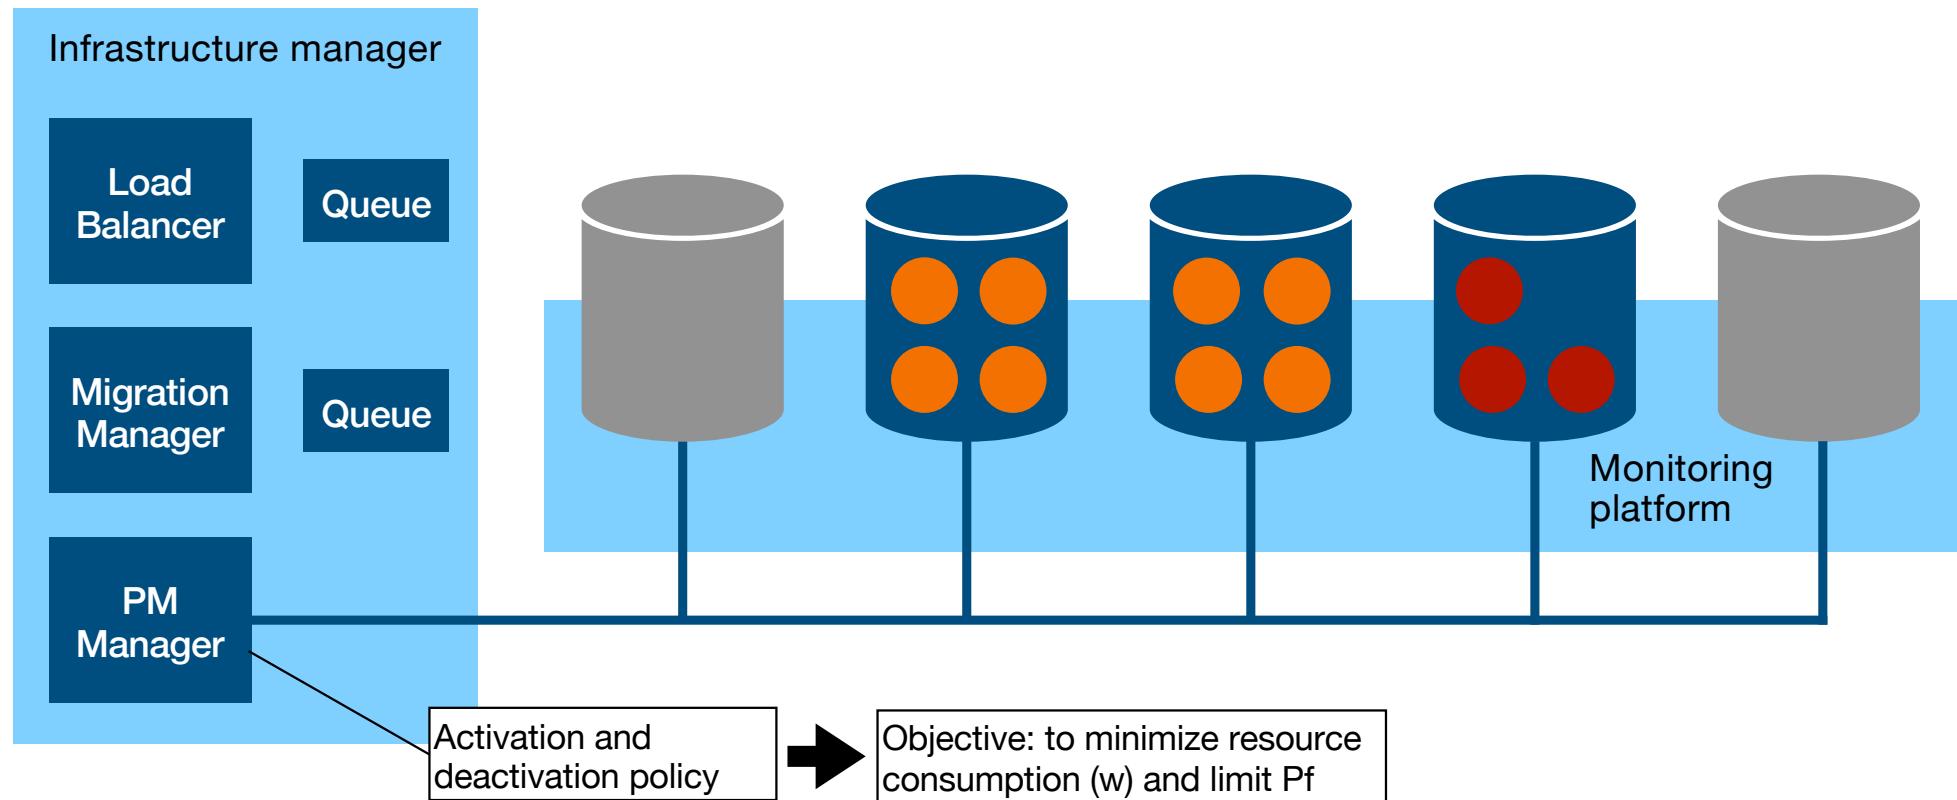
# Example



# Example



# Example



# Static optimization

# System

## Modeling hypotheses

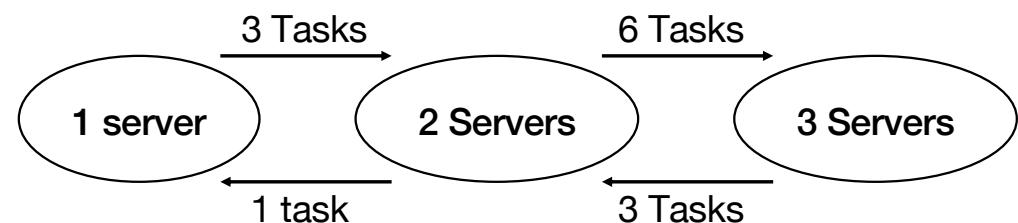
- Tasks arrive following Poisson  $\lambda$  and exponential service times  $\mu^{-1}$
- A server can serve up to  $N$  tasks, there are  $M$  servers
- Exponential boot up times:  $\alpha^{-1}$  Immediate shut downs.
- Server lifetimes are exponential [1]  $\nu^{-1}$
- Energy consumption: 0 off,  $P_{\text{idle}}$  when on,  $P_{\text{load}}$  proportional to load [2]
- Tasks can be moved between servers before crashing

1. K. S. Trivedi y A. Bobbio, «Reliability and Availability Engineering: Modeling, Analysis, and Applications», Cambridge University Press, 2017
2. Gong Chen et al. «Energy-Aware Server Provisioning and Load Dispatching for Connection-Intensive Internet Services». NSDI'08.

# System

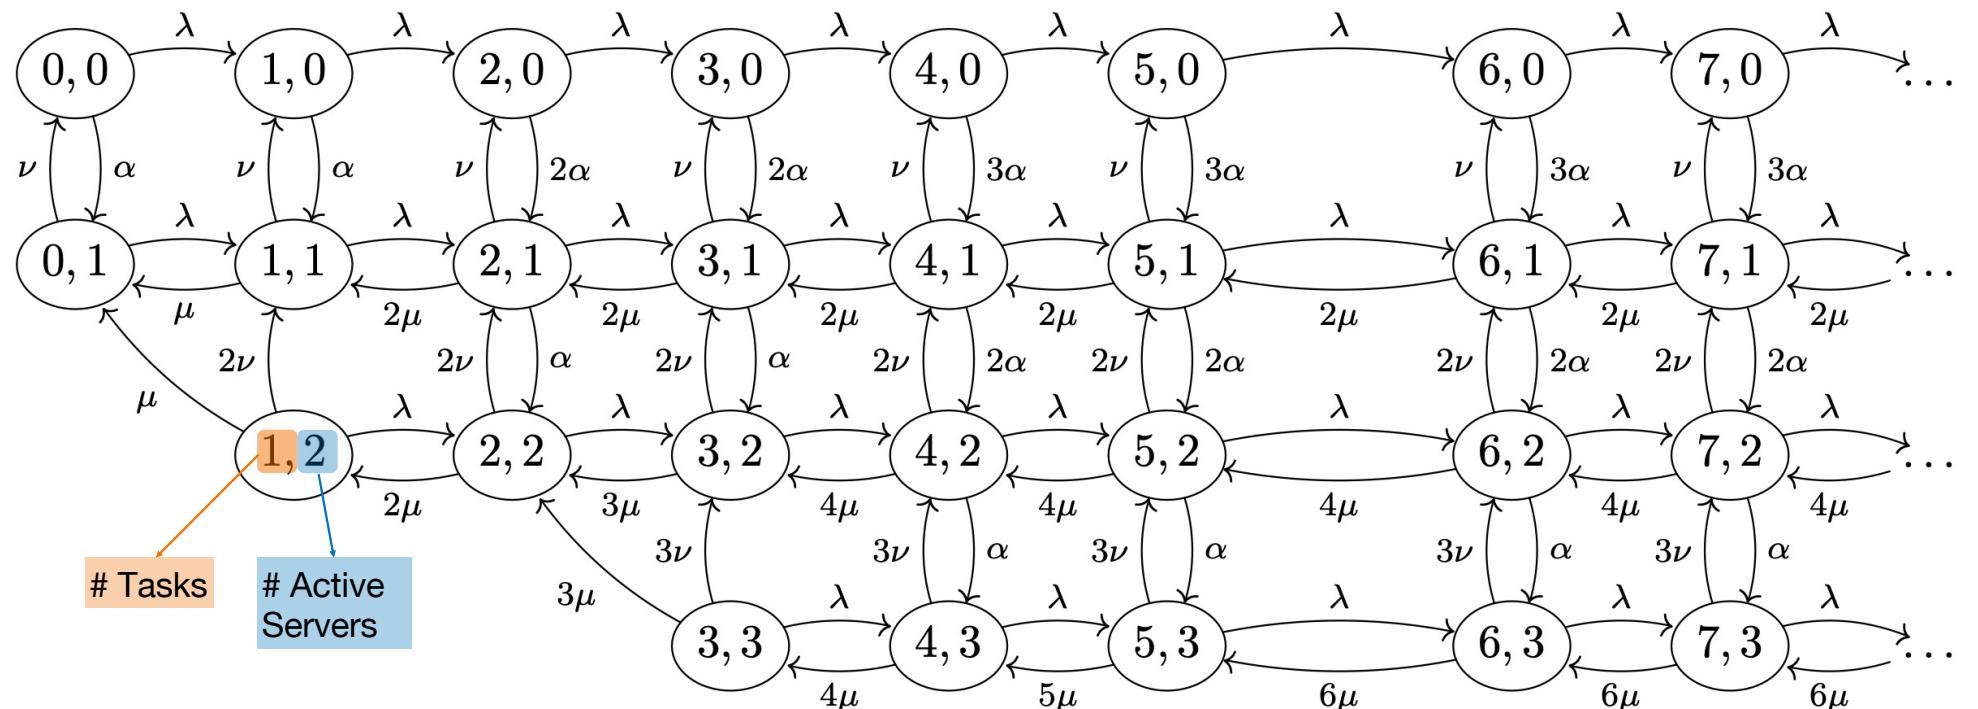
## De/Activation Policy

- Activation threshold  $t_{on}(m)$  : Number of tasks on the system that initiates activation of the  $m$ -th server
- Deactivation threshold  $t_{off}(m)$  : Number of tasks on the system that causes a server to be disabled when there are  $m$  active servers
- Example
  - $t_{on}(2) = 3, t_{on}(3) = 6$
  - $t_{off}(2) = 1, t_{off}(3) = 3$



# Analytical model: Markov chain

## Quasi-Birth-Death process

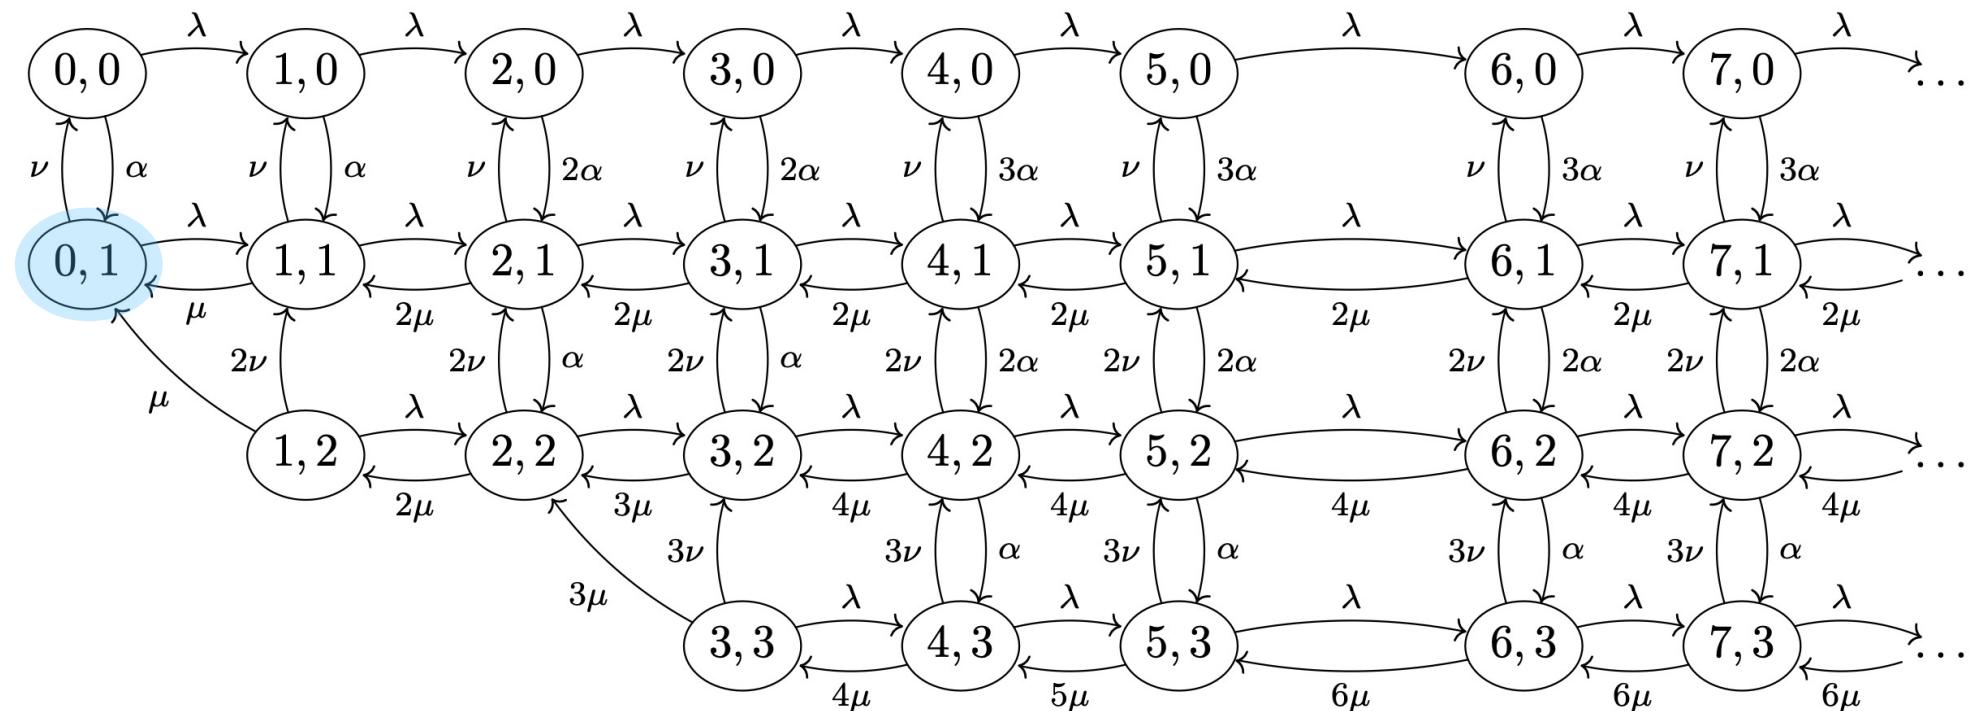


Niveles iniciales

Niveles repetitivos

# Analytical model: Markov chain

## Quasi-Birth-Death process

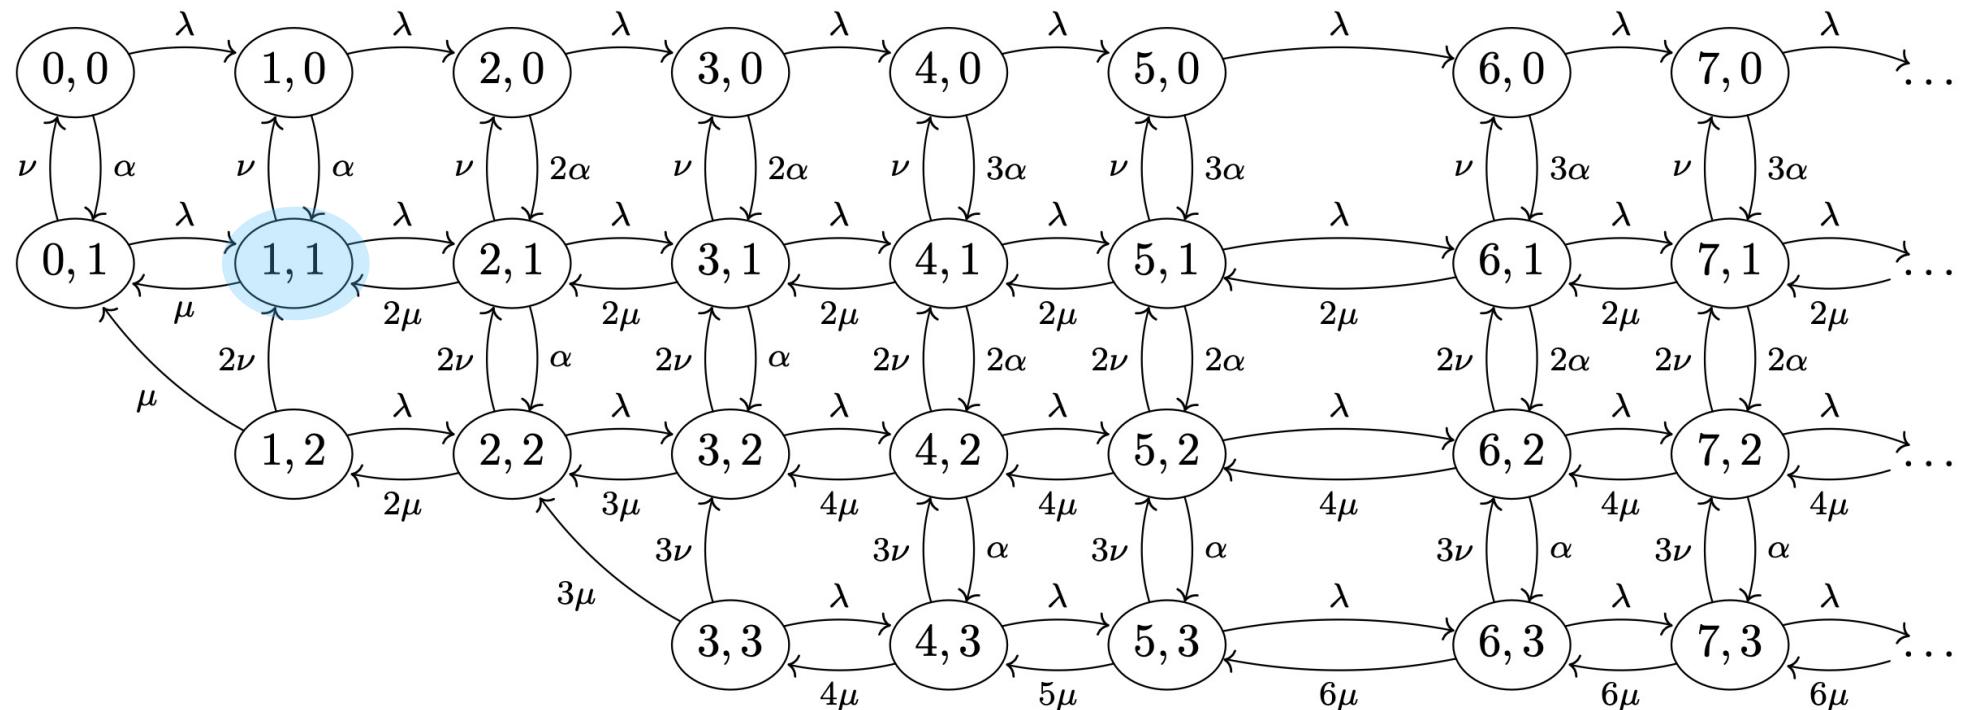


Niveles iniciales

Niveles repetitivos

# Analytical model: Markov chain

## Quasi-Birth-Death process

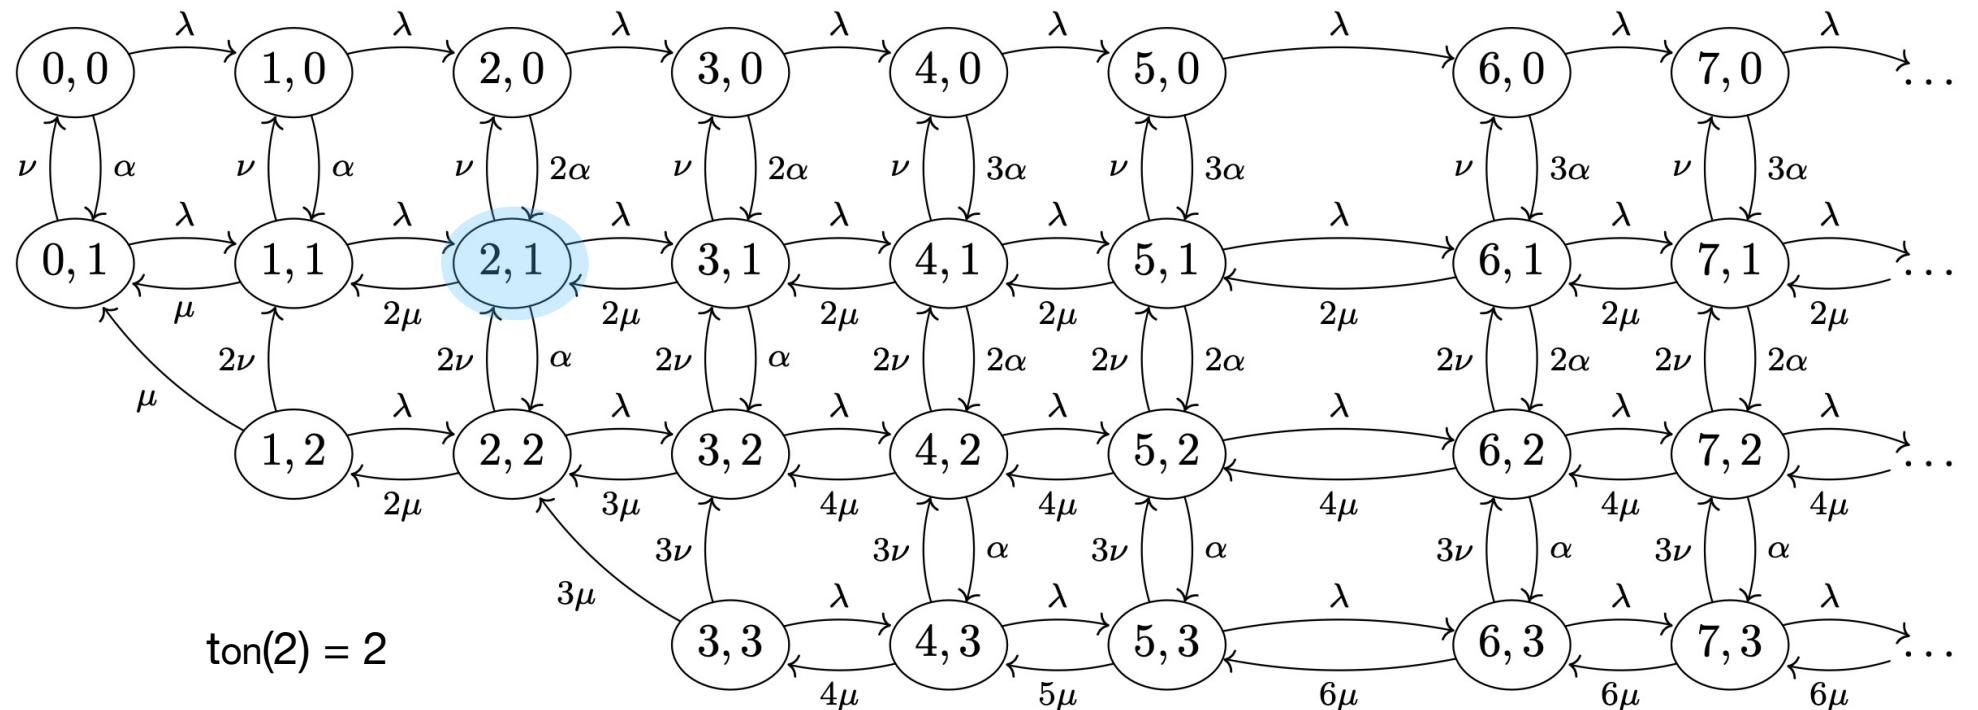


Niveles iniciales

Niveles repetitivos

# Analytical model: Markov chain

## Quasi-Birth-Death process

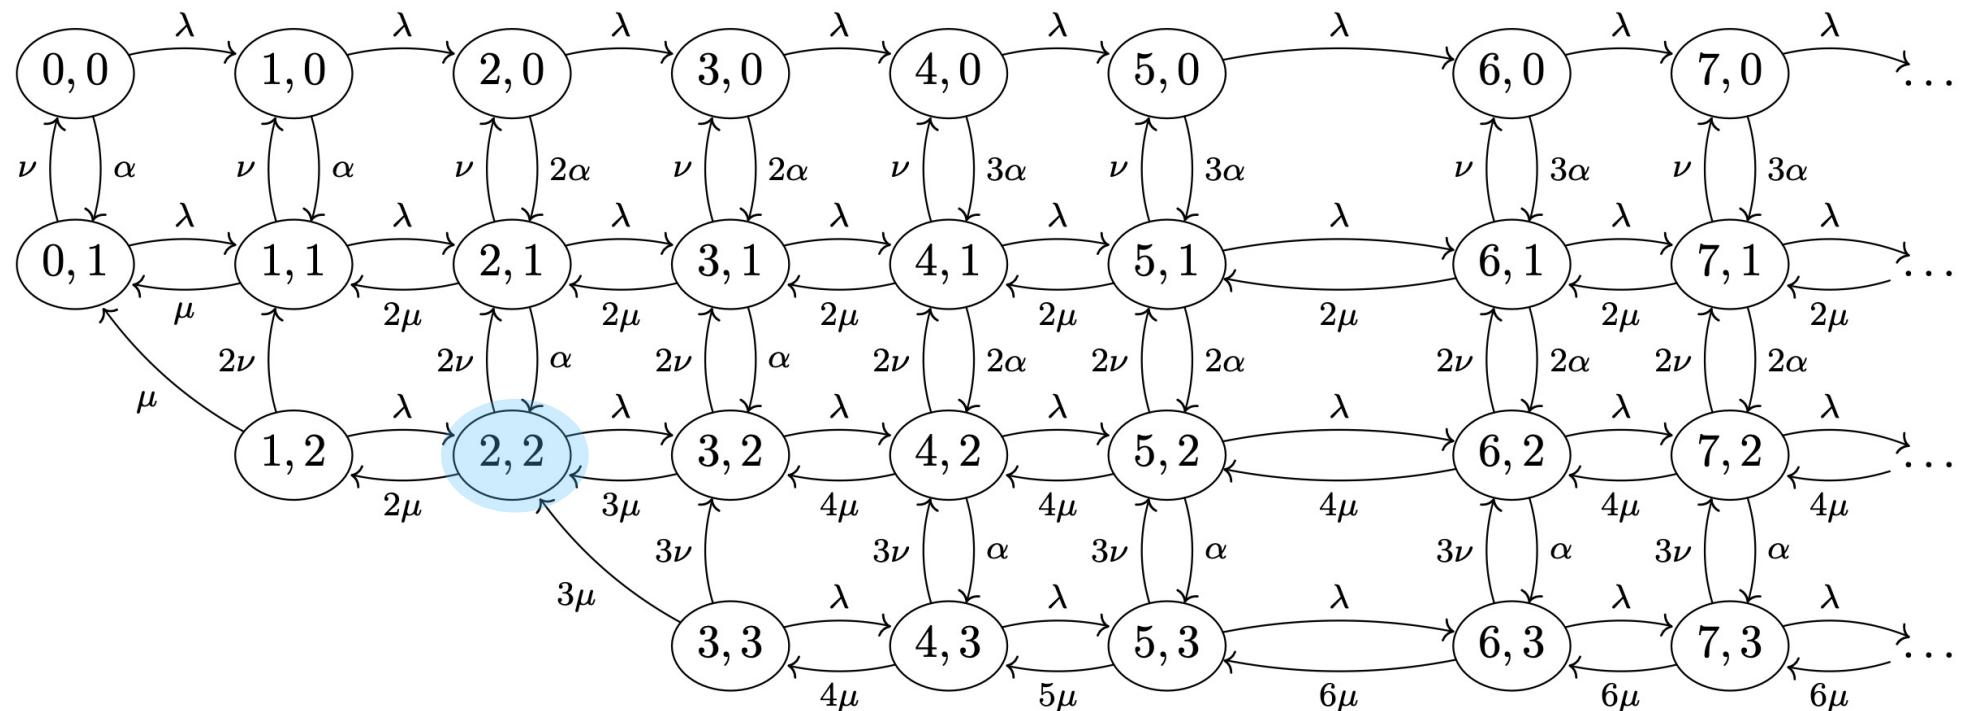


Niveles iniciales

Niveles repetitivos

# Analytical model: Markov chain

## Quasi-Birth-Death process

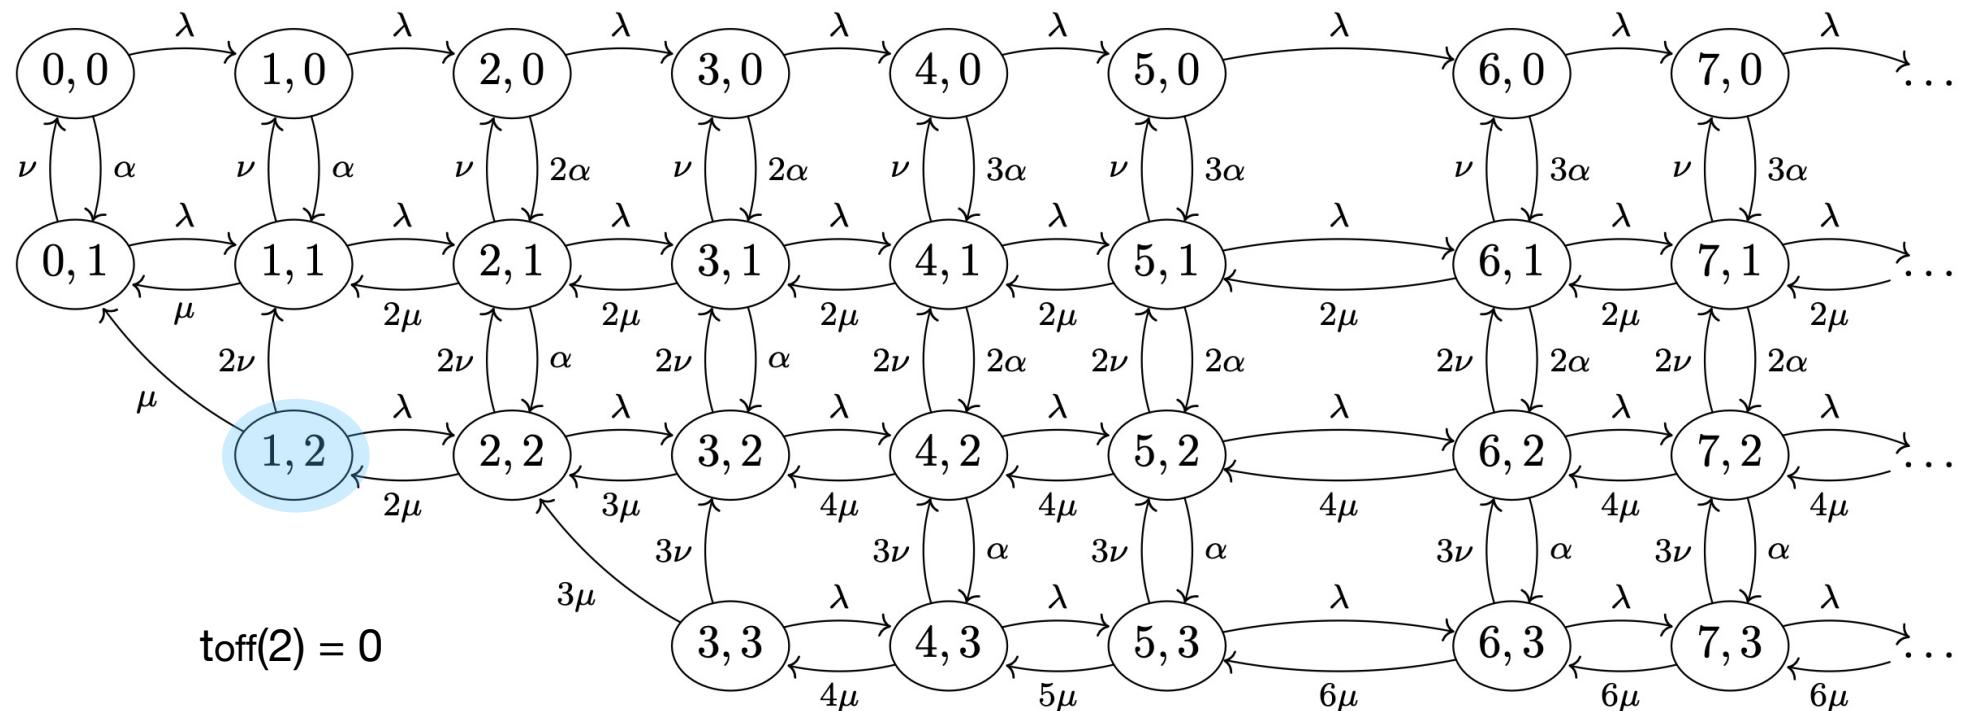


Niveles iniciales

Niveles repetitivos

# Analytical model: Markov chain

## Quasi-Birth-Death process

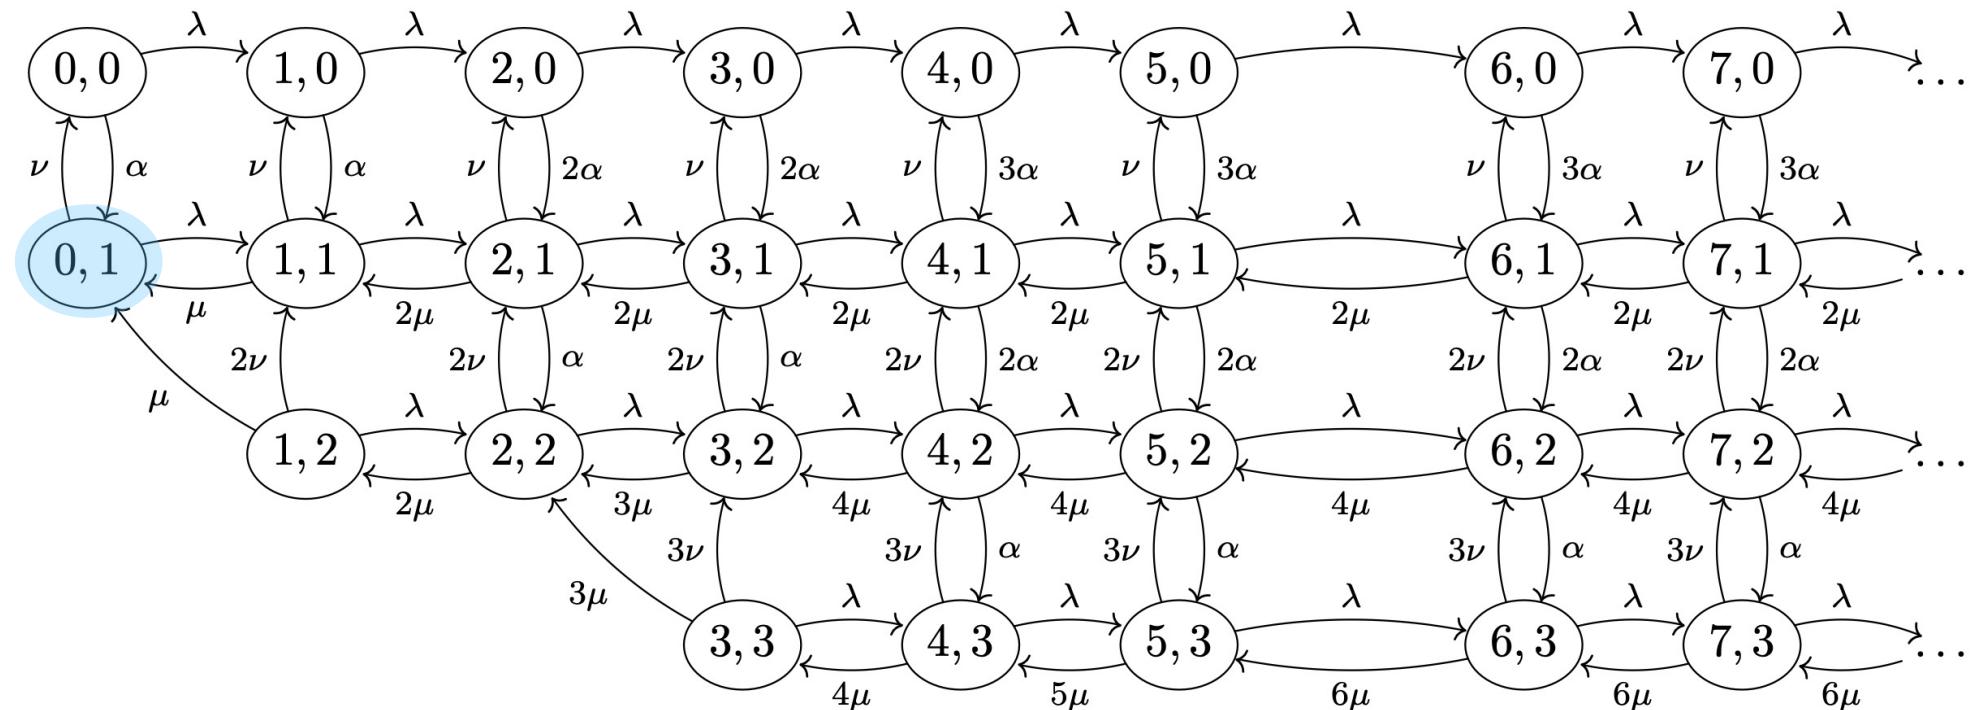


Niveles iniciales

Niveles repetitivos

# Analytical model: Markov chain

## Quasi-Birth-Death process

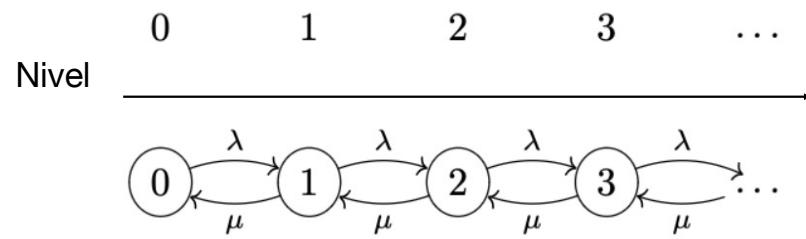


## Niveles iniciales

## Niveles repetitivos

# Birth-Death vs. Quasi-Birth-Death

## Birth death process



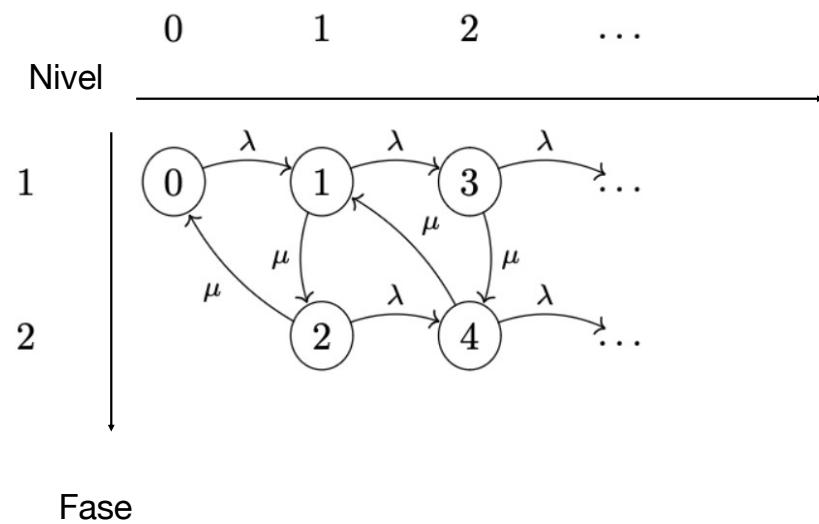
$$Q = \begin{pmatrix} -f_0 & f_0 & & & & \\ b_1 & -(b_1 + f_1) & f_1 & & & \\ & b_2 & -(b_2 + f_2) & f_2 & & \\ & & & \ddots & & \\ & & & & \ddots & \ddots \end{pmatrix}$$

$$f_i = \lambda$$

$$b_j = \mu$$

# Birth-Death vs. Quasi-Birth-Death

## Quasi-birth-death process



$$\mathbf{F}^{(0)} = (\lambda \quad 0)$$

$$\mathbf{B}^{(1)} = \begin{pmatrix} 0 \\ \mu \end{pmatrix} \quad \mathbf{L}^{(0)} = -\lambda$$

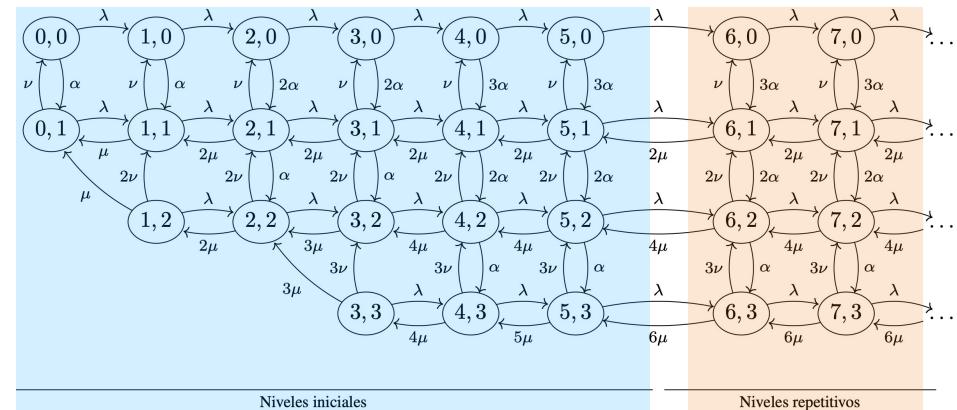
$$\mathbf{Q} = \begin{pmatrix} \mathbf{L}^{(0)} & \mathbf{F}^{(0)} & & & \\ \mathbf{B}^{(1)} & \mathbf{L}^{(1)} & \mathbf{F}^{(1)} & & \\ & \mathbf{B}^{(2)} & \mathbf{L}^{(2)} & \mathbf{F}^{(2)} & \dots \\ & & & \vdots & \ddots \end{pmatrix}$$

$$\mathbf{F}^{(l)} = \begin{pmatrix} \lambda & 0 \\ 0 & \lambda \end{pmatrix} \quad \mathbf{L}^{(l)} = \begin{pmatrix} -(\mu + \lambda) & \mu \\ 0 & -(\mu + \lambda) \end{pmatrix}$$

$$\mathbf{B}^{(l)} = \begin{pmatrix} 0 & 0 \\ \mu & 0 \end{pmatrix}$$

# Markov chain

## Various initial levels



$$\mathbf{Q} = \left[ \begin{array}{cccccc}
 B_{00} & B_{01} & 0 & 0 & \dots & 0 \\
 B_{10} & B_{11} & B_{12} & 0 & \dots & 0 \\
 0 & B_{21} & B_{22} & B_{23} & \dots & 0 \\
 \vdots & \vdots & \vdots & \vdots & \ddots & \dots \\
 0 & \dots & \dots & B_{I-2,I-3} & B_{I-2,I-2} & B_{I-2,I-1} \\
 0 & \dots & \dots & 0 & B_{I-1,I-2} & B_{I-1,I-1} \\
 0 & \dots & \dots & 0 & 0 & A_2 \\
 \vdots & \vdots & \vdots & \vdots & \ddots & \dots
 \end{array} \right] \quad \mathbf{B}: \text{matriz para los niveles iniciales}$$

# Theorem

## Convergence Assurance

$\pi^A$  **Theorem 1. The QBD process is stable if it holds**

$$\frac{\lambda}{NM\mu} < \frac{\alpha}{\nu + \alpha}$$

- Proof: the drift of the system to the higher levels has to be less than the drift to the lower levels [1]

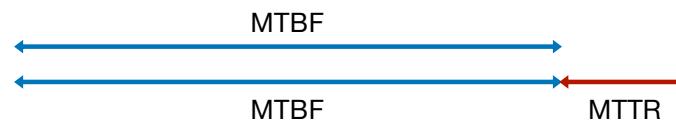
$$\pi_A \mathbf{A}_0 \mathbf{1} < \pi_A \mathbf{A}_2 \mathbf{1},$$

- Distribution Vector

$$\mathbf{A} = \mathbf{A}_0 + \mathbf{A}_1 + \mathbf{A}_2$$

Interpretation:

$$\rho < \frac{\text{MTBF}}{\text{MTTR} + \text{MTBF}}$$



- Marcel F. Neuts, «Matrix-Geometric Solutions in Stochastic Models: An Algorithmic Approach», Dover Publications, 1995. ISBN: 978-0486683423

# Metrics

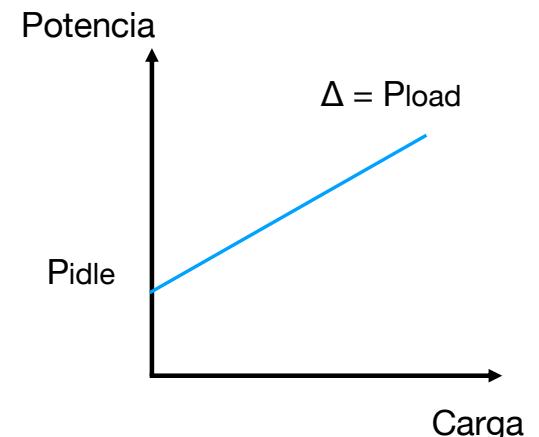
## Power Consumption ( $w$ )

- Markov chain solved

$\pi_{i,j} \rightarrow$  Prob. de  $i$  tareas y  $j$  servidores

- Average power consumed

$$\omega = \sum_{(i,j) \in S} \pi_{i,j} j \left( P_{\text{idle}} + P_{\text{load}} \frac{\min(i, jN)}{jN} \right)$$



# Metrics

## Probability of failure (Pf)

- Assuming Probabilities  $\approx 0$

$$P_f = P_{\text{wait}} + P_{\text{int}}.$$

- Failure bc wait on arrival (PASTA)

$$P_{\text{wait}} = \sum_{(i,j) \in S} \mathbf{1}(i \geq jN) \pi_{i,j},$$

Task drop rate

- Failure bc server crash

$$P_{\text{int}} = \frac{\gamma}{\lambda}$$

$$\gamma = \sum_{(i,j) \in S} \pi_{i,j} j \nu F_{i,j}^r,$$

Indicator function

Server crash rate

Number of Tasks Affected

Capacity after crash

$$F_{i,j}^r = \min(\max(i - (j-1)N, 0), N),$$

# Validation

## Simulation parameters

Rack

8 Servers Dell Power Edge 32 GB  
 Pmax 270 W y Pidle 150 W [1]  
 Boot up: 3 min [measured]  
 MTBF: 768 hours [2]

Nano

64 Servers Raspberry Pi 4b 4 GB  
 Pmax 7.6 W y Pidle 4.6 W [3]  
 Boot up: 20 s [measured]  
 MTBF: 1/4 of the above

Párametro	Servidores rack	Servidores nano
M	8 servidores	64 servidores
N	32 tareas/servidor	4 tareas/servidor
$1/\mu$	1 hora	1 hora
$1/\alpha$	3 min.	20 s.
$1/\nu$	32 días	8 días
$P_{idle}$	150 W	4.6 W
$P_{load}$	120 W	3.0 W

$$1/\mu = 1 \text{ h}$$

1. TPCDB: <http://www.tpcdb.com/product.php?id=2325>
2. G. L. Santos et al., «Analyzing the IT subsystem failure impact on availability of cloud services». ISCC 2017
3. TPCDB: <http://www.tpcdb.com/product.php?id=4417>

# Validation

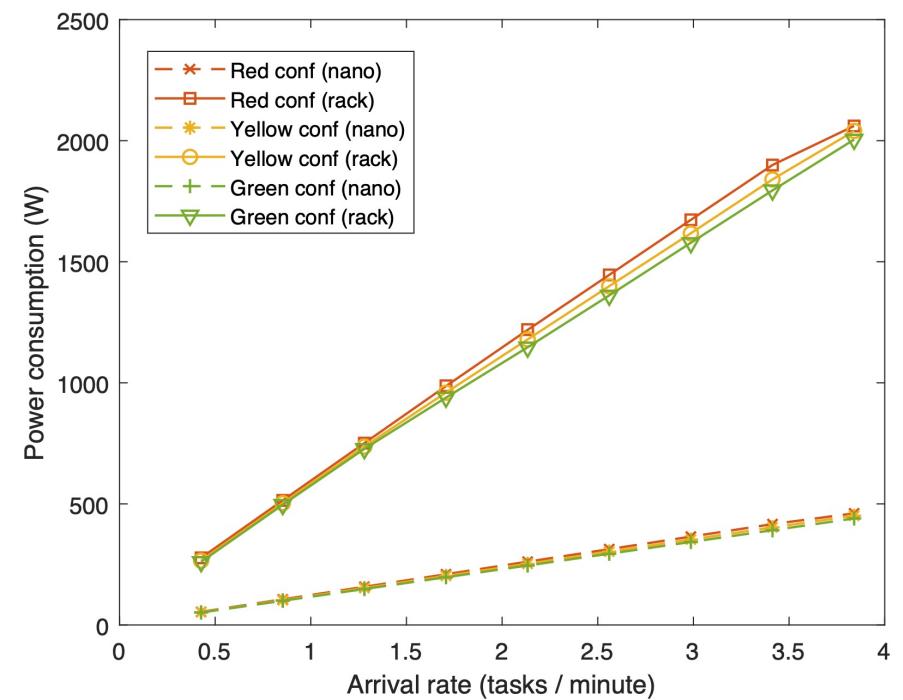
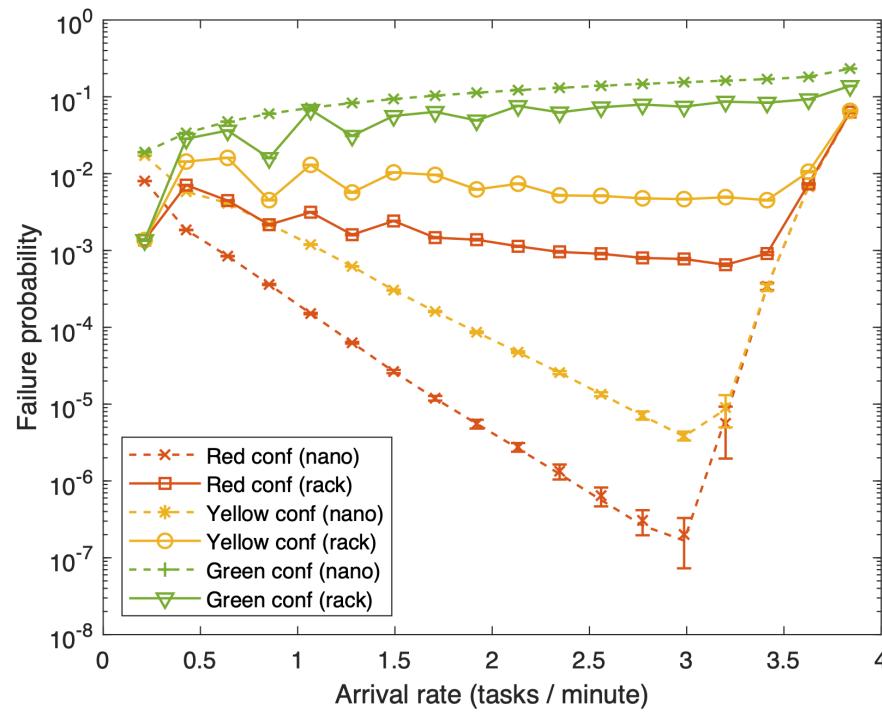
## Power On and Off Policies

Example  
Serv. rack (M=32)

<b>Green</b>	Turn on if all active servers are busy Turn off if, doing so, leaves room for a task	$ton(5) = 128$ $toff(5) = 127$
<b>Red</b>	Turn on when occupancy exceeds 95% Turn off if occupancy would fall below 85%	$ton(5) = 122$ $toff(5) = 108$
<b>Yellow</b>	Turn on when occupancy exceeds 95% Turn off if occupancy would fall below 95%	$ton(5) = 122$ $toff(5) = 121$

# Results

## Failure probability & Power consumption



# Optimal configuration

## Formulation

- Model allows predicting performance for a given configuration
- Challenge: meeting a reliability criterion (prob. failure) and minimizing consumption

$$\min_{\{t_m^{\text{on}}\}, \{t_m^{\text{off}}\}} \omega$$

Bound for  $P_f$   
*Target failure probability*

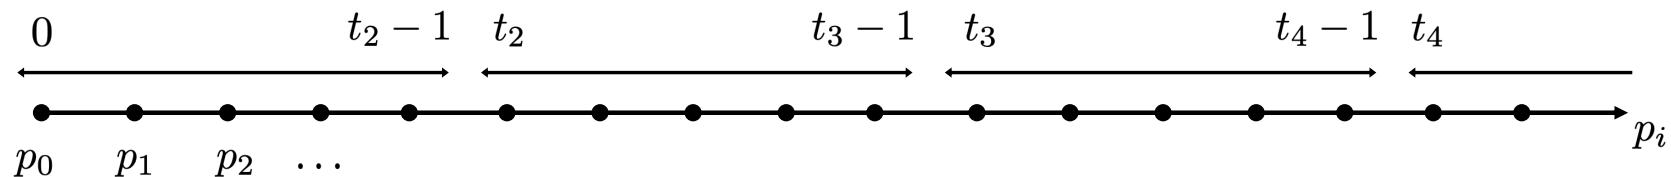
$$\text{subject to } P_f \leq T_f$$

- Problem: Large space size of possible combinations

# Optimal configuration (approach)

## Assumptions

- Threshold policy:  $t_m^{\text{off}} = t_m^{\text{on}} - 1 \longrightarrow t_m^{\text{on}} \equiv t_m$
- Server crash and activation does not affect task distribution  $\{p_i\}$ 
  - Behaves as a classical M/M/C,  $C = M \times N$
  - Power consumed as a function of thresholds
  - Decoupling the effect of each threshold on Pf:



$$\omega_{m,t_m}^{\text{idle}} = \sum_{i=t_m}^{\infty} p_i P_{\text{idle}}$$

$$\omega = \sum_{m=1}^M \omega_{m,t_m}^{\text{idle}} + \sum_{i=1}^{\infty} p_i \min(i, C) P_{\text{req}}$$

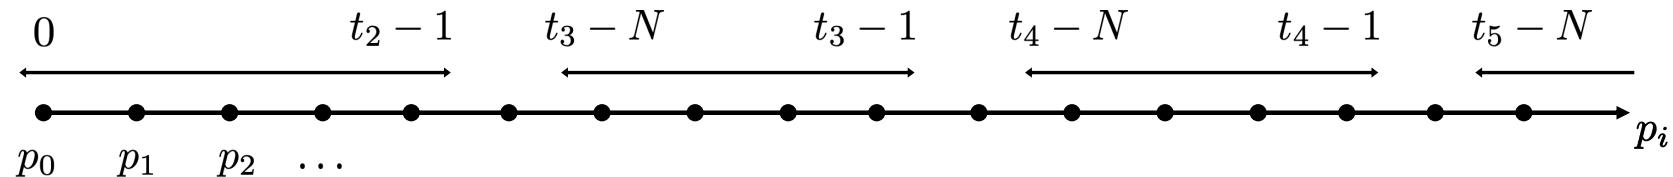
# Optimal configuration (approach)

## Assumptions

- Threshold policy:  $t_m^{\text{off}} = t_m^{\text{on}} - 1 \longrightarrow t_m^{\text{on}} \equiv t_m$
- Server crash and activation does not affect task distribution  $\{p_i\}$ 
  - Behaves as a classical M/M/C,  $C = M \times N$
  - Power consumed as a function of thresholds
  - Decoupling the effect of each threshold on  $P_f$ :

$$\omega_{m,t_m}^{\text{idle}} = \sum_{i=t_m}^{\infty} p_i P_{\text{idle}}$$

$$\omega = \sum_{m=1}^M \omega_{m,t_m}^{\text{idle}} + \sum_{i=1}^{\infty} p_i \min(i, C) P_{\text{req}}$$



# Computing the optimal configuration

## *Multiple-choice knapsack problem*

$$\begin{array}{ll}
 \min_{\{t_m^{\text{on}}\}, \{t_m^{\text{off}}\}} & \omega \\
 \text{sujeto a} & P_f \leq T_f \\
 \end{array}
 \quad \Rightarrow \quad
 \begin{array}{ll}
 \min_{\{x_{m,k}\}} & \sum_{m=2}^M \sum_{k \in \mathcal{M}_m} x_{m,k} \omega_{m,k}^{\text{idle}} \\
 \text{sujeto a} & \sum_{m=2}^M \sum_{k \in \mathcal{M}_m} x_{m,k} P_{m,k}^f \leq T_f \\
 & \sum_{k \in \mathcal{M}_m} x_{m,k} = 1, \quad \forall m = 2, \dots, M \\
 & x_{m,k} \in \{0, 1\}, \quad \forall m = 2, \dots, M, k \in \mathcal{M}_m
 \end{array}$$

1 if the threshold 'm' is equal to 'k'

### Post - optimization

Threshold adjustment  $t_m$

If  $P_f > T_f$ , lower it

If  $P_f < T_f$ , increase it

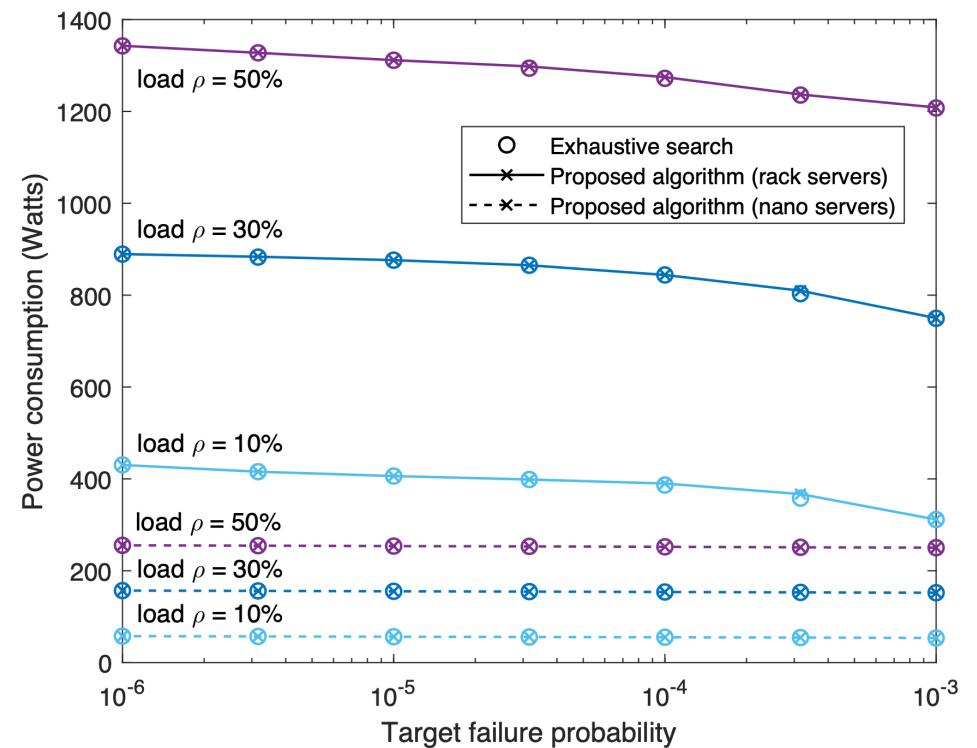
# Evaluation

## Validation of the proposed configuration

- Different load and Tf values
- Complexity
  - Algorithm: 20 s -- 37 s
  - Search: 3 h (rack), 10 d (nano)
- Nano servers are more efficient
 

↓

Pseudo exhaustive



# Evaluation

## Comparison with other strategies

Heuristic  
[1]

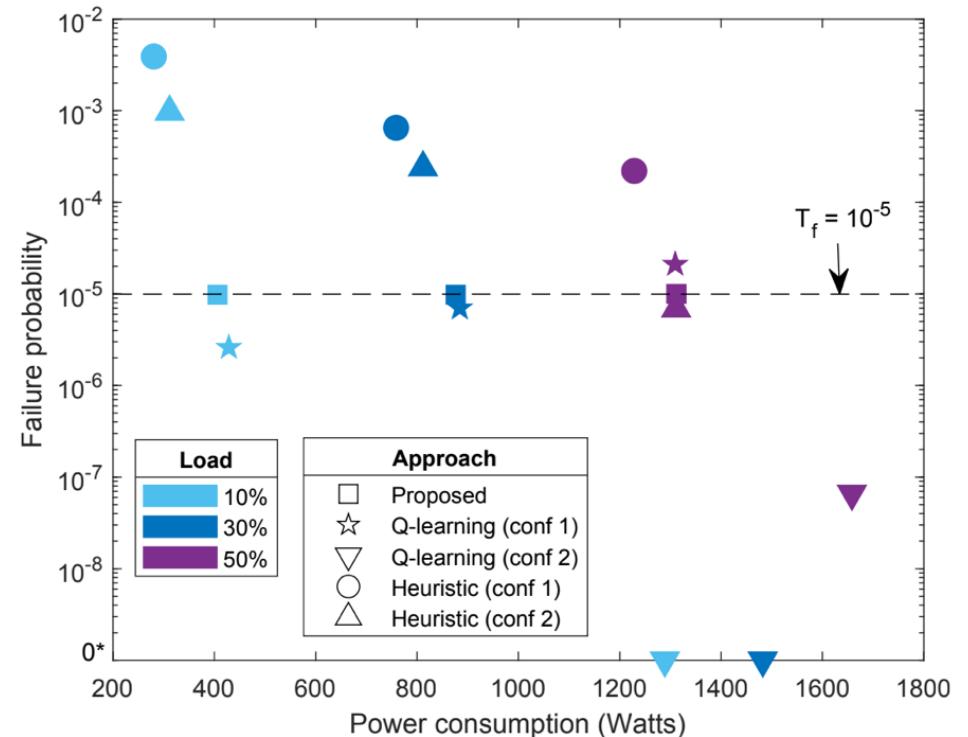
Two threshold sets to de/activate  
*Minimization of migrations* (MM)  
Threshold: {50%, 90%} y {40%, 80%}

Q-learning  
[2]

State: # tasks, # servers  
Actions: Power on/off/keep  
Penalty: failures y consumption

$$p_t = \beta n_t + (1 - \beta) \omega_t, \\ \beta = \{0.8, 1\}$$

1. A. Beloglazov et al. «Energy-aware resource allocation heuristics for efficient management of data centers for cloud computing», Future generation computer systems 28.5 (2012), págs. 755-768.
2. S. Telenyk et al. «Modeling of the Data Center Resource Management Using Reinforcement Learning», PIC S&T 2018



# Dynamic optimization

# Motivation

## Static vs. Dynamic

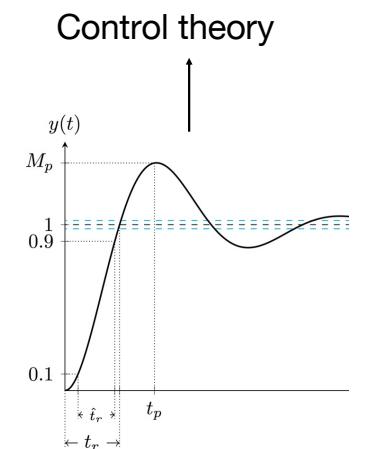
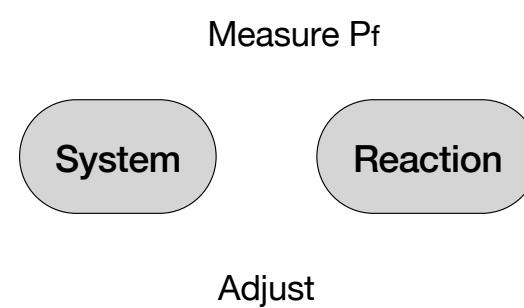
- Static optimization: calculation of thresholds a priori

- Need to estimate parameters
- Based on certain modeling hypotheses

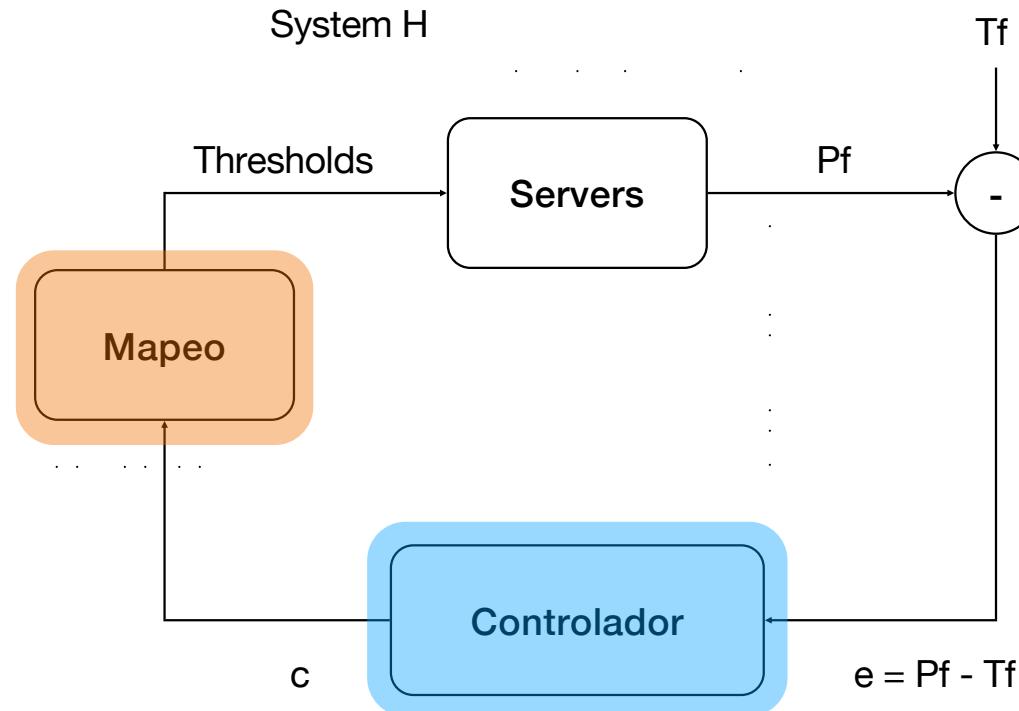
Poisson Arrivals  
Exponential service  
Exponential boots

- Dynamic optimization: changing thresholds as a reaction

- Lower thresholds  $\longrightarrow \downarrow P_f \uparrow \omega$
- Raise thresholds  $\longrightarrow \uparrow P_f \downarrow \omega$



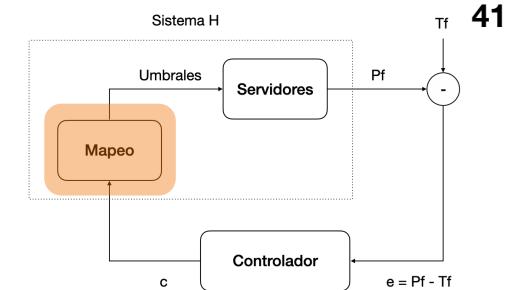
# System Design: A3S



# System Design: A3S

## Mapping: from 'c' to thresholds

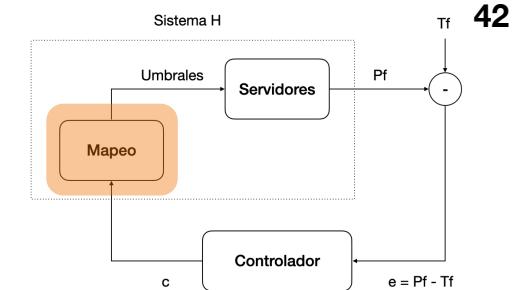
- Threshold Policy  $t_m^{\text{off}} = t_m^{\text{on}} - 1 \longrightarrow t_m^{\text{on}} \equiv t_m$
- Error  $e = P_f - T_f$ 
  - Positive Error -> Lower Thresholds
  - Negative Error -> Raising Thresholds
- Minimum thresholds: all on
- Maximum thresholds: only turns on when full **(green)**



# System Design: A3S

## Mapping: from 'c' to thresholds

- Threshold Policy  $t_m^{\text{off}} = t_m^{\text{on}} - 1 \longrightarrow t_m^{\text{on}} \equiv t_m$
- Error  $e = P_f - Tf$ 
  - Positive Error  $\rightarrow$  Lower Thresholds
  - Negative Error  $\rightarrow$  Raising Thresholds
- Minimum thresholds: all on
- Maximum thresholds: only turns on when full **(green)**



42

Example  
M=4, N=3

c	t <sub>1</sub>	t <sub>2</sub>	t <sub>3</sub>	t <sub>4</sub>
0	0	3	6	9
1	0	<b>2</b>	6	9
2	0	2	<b>5</b>	9
3	0	2	5	<b>8</b>
4	0	<b>1</b>	5	8
...	...	...	...	...
18	0	0	0	0

$$q = \left\lfloor \frac{c}{M-1} \right\rfloor, \quad r = c - (M-1) \left\lfloor \frac{c}{M-1} \right\rfloor$$

$$t_m = \begin{cases} (m-1)M - q - 1, & \text{si } m \leq r + 1, \\ (m-1)M - q, & \text{en otro caso.} \end{cases}$$

# System Design: A3S

## Controller: Proportional Integral (PI)

- Control signal

$$c(t) = K_p e(t-1) + K_i \sum_{t'=0}^{t-2} e(t')$$

- Parameter values (Ziegler–Nichols) [1, 2]

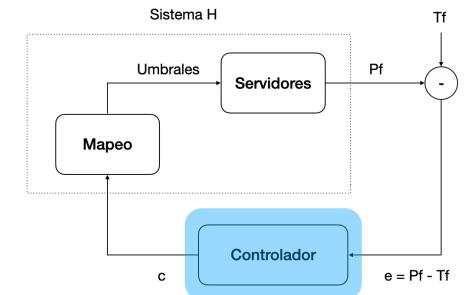
$$K_p = \frac{0.4}{\hat{H}}$$

$$K_i = \frac{0.4}{\hat{H} \cdot 0.85 \cdot 2}$$

donde  $\hat{H} > |H|$

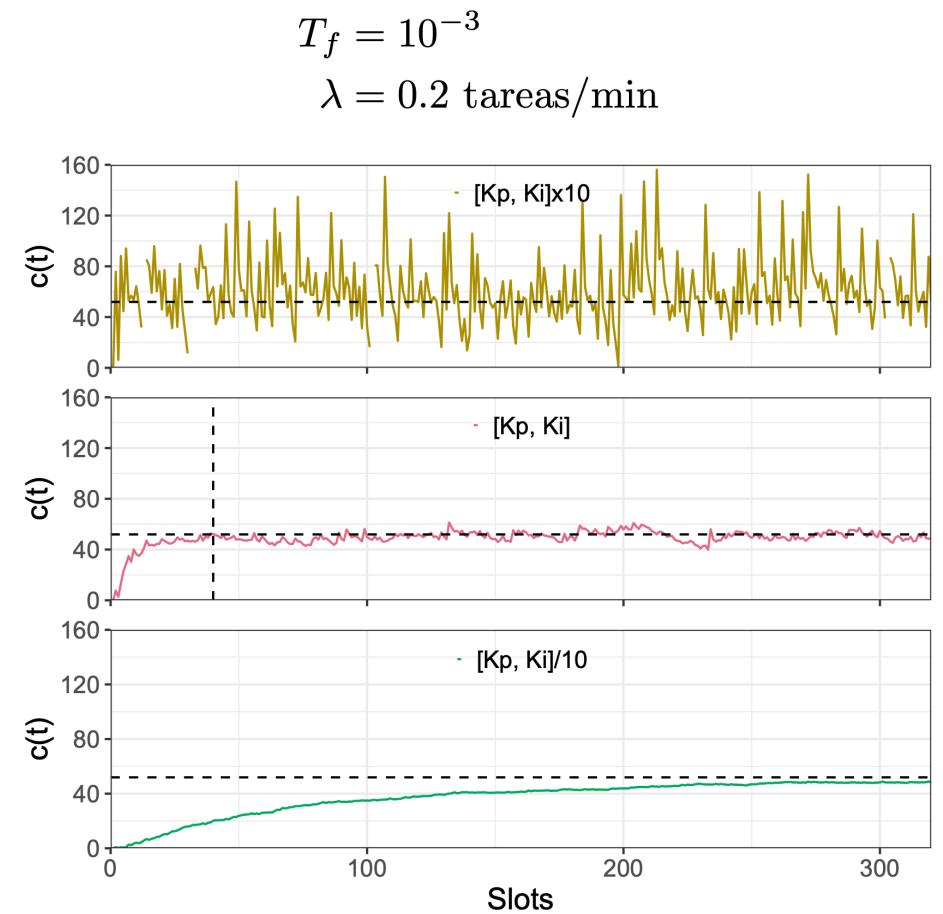
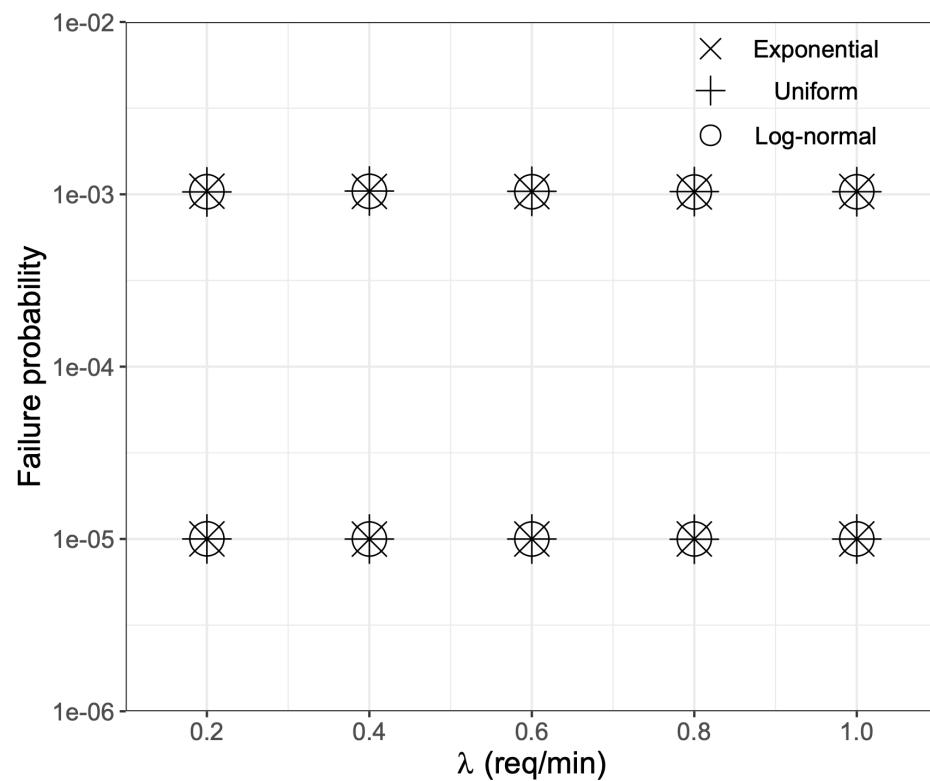
$$\hat{H} = \frac{T_f}{M - 1}$$

1. A. Garcia-Saavedra et al., «Adaptive Mechanism for Distributed Opportunistic Scheduling», IEEE Transactions on Wireless Communications, 2015
2. P. Serrano et al., «Control Theoretic Optimization of 802.11 WLANs: Implementation and Experimental Evaluation», Elsevier Computer Networks, 2013



# Validation

## Efficiency and Configuration

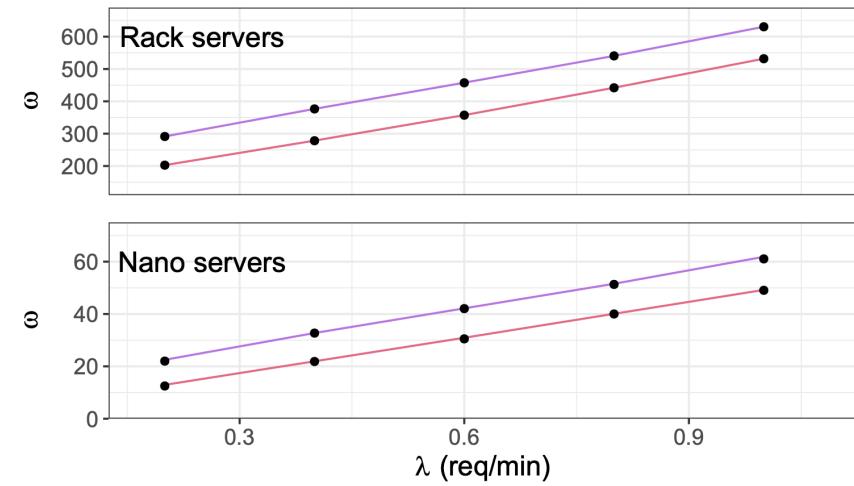


# Evaluation

## Vs. Search & Static Optimal

- Comprehensive threshold search
- Vs. Static configuration
  - Three 24-hour periods of Google Workflow trace [1]
  - $M \times N = 128$
  - $I = \{30, 44, 64\}$  Erlangs
  - $T_f = \{10^{-3}, 10^{-5}\}$

1. Google, «Workflow Trace Archive Google trace», Zenodo, Jun. 24, 2019. doi: 10.5281/zenodo.3254540.



Load	A3S		Static	
	$P_f$	$\omega$ (W)	$P_f$	$\omega$ (W)
Small	$10^{-3}$	516.1	$0.3 \cdot 10^{-3}$	665.2
	$10^{-5}$	687.4	$0.8 \cdot 10^{-5}$	710.8
Medium	$10^{-3}$	727.5	$0.4 \cdot 10^{-3}$	773.5
	$10^{-5}$	880.8	$0.4 \cdot 10^{-5}$	942.3
High	$10^{-3}$	1 020.3	$0.5 \cdot 10^{-3}$	1 062.1
	$10^{-5}$	1 212.8	$0.8 \cdot 10^{-5}$	1 232.5

# Comparison

## Vs. Reinforcement learning (RL)

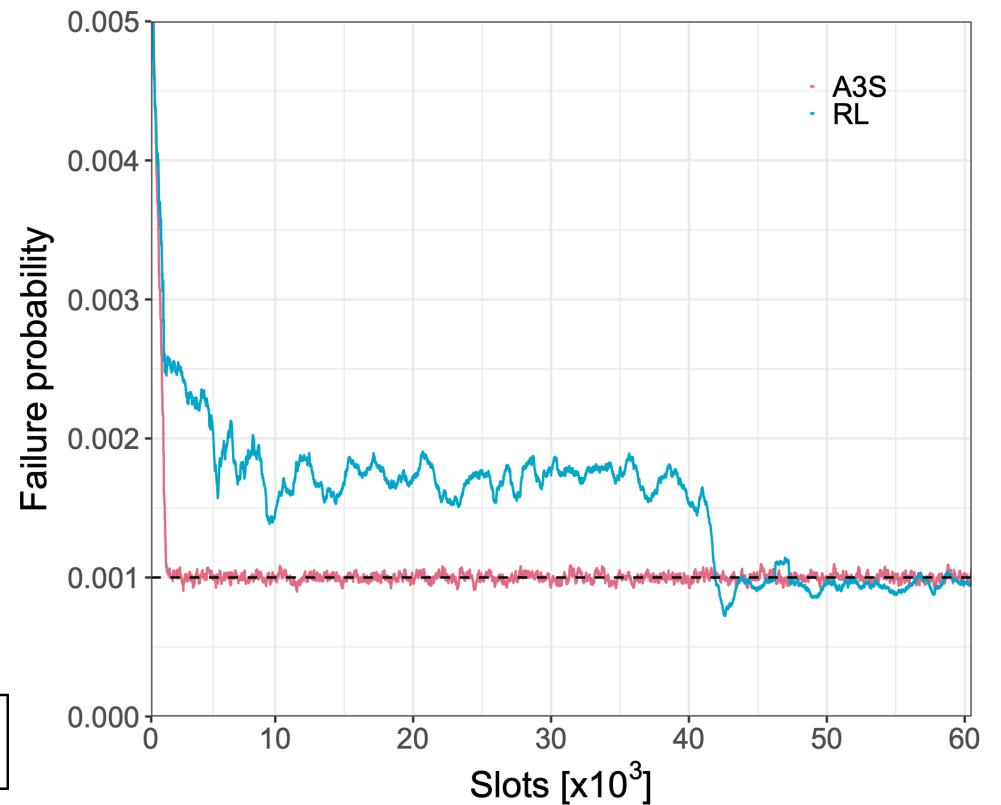
1. S. Telenyk et al. «Modeling of the Data Center Resource Management Using Reinforcement Learning», PIC S&T 2018

- Same technique as before [1]
- Status: # tasks, # servers
- Action: Turn on/off/hold
- Numerical search to achieve  $T_f$

$$p_t = \beta n_t + (1 - \beta) \omega_t$$

- Convergence: 90% Qvalues

$$Q(s_t, a_t) = (1 - \alpha)Q(s_t, a_t) + \alpha \left[ p_t + \gamma \min_{a_{t+1}} Q(s_{t+1}, a_{t+1}) \right]$$



# Conclusiones y Trabajo futuro

## Conclusiones

- The softwarization of networks poses opportunities and challenges
- Sustainable operation but guaranteeing high performance
- Two solutions
- Static: Estimation
- Dynamic: adaptation
- Advantages vs. RL

## Trabajo futuro

- Other service models: heterogeneous requests
- Optimizing the Farm Design
- Support for different services with different requirements
- Impact of the transmission medium

# Additional information

## Publications

### Static optimization

- J. Ortín et al., «Analysis of scaling policies for NFV providing 5G/6G reliability levels with fallible servers», **IEEE Transactions on Network and Service Management (JCR Q2)**, Junio 2022

### Dynamic optimization

- J. Perez-Valero et al., «Energy-Aware Adaptive Scaling of Server Farms for NFV with Reliability Requirements», **IEEE Transactions on Mobile Computing (JCR Q1)**, Junio 2023

### Other scenarios

- J. Perez-Valero et al., «Performance Trade-offs of Auto Scaling Schemes for NFV with Reliability Requirements», **Computer Communications (JCR Q1)**, Diciembre 2023

### Design optimization

- J. Pérez-Valero et al. «Minimum-Cost Design of Auto-Scaling Server Farms Providing Reliability Guarantees», **IEEE Open Journal of the Communications Society (JCR Q1)**, Julio 2025